From Bayesian Inference to Neural Computation: The Analytical Emergence of Neural Network Structure from Probabilistic Relevance Estimation

Jaepil Jeong
Cognica, Inc.
Email: jaepil@cognica.io
Date: February 13, 2026

"The theory of probabilities is at bottom nothing but common sense reduced to calculus."
— Pierre-Simon Laplace, *Théorie analytique des probabilités*, 1812

Abstract

We demonstrate that feedforward neural network structure emerges analytically from first-principles Bayesian inference over multiple relevance signals in information retrieval. Starting from a single question — *what is the probability that a document is relevant given multiple evidence signals?* — we apply Bayes' theorem to derive sigmoid calibration of individual scores, introduce a log-odds conjunction framework that resolves the well-known shrinkage problem of naive probabilistic conjunction, and show that the resulting end-to-end computation has the structure of a feedforward neural network: inputs are calibrated to probabilities, pass through a logit nonlinearity into log-odds space (the natural parameter space of Bernoulli random variables), undergo linear aggregation (normalized Logarithmic Opinion Pooling with confidence scaling), and pass through a sigmoid activation (posterior computation). When signals share sigmoid calibration, the logit-sigmoid identity collapses the hidden layer and the computation reduces to logistic regression — confirming the classical Bayesian-logistic correspondence from a new direction. When signals have heterogeneous calibrations (the practically important case of hybrid search combining BM25 and vector similarity), the logit performs a genuine nonlinear transformation, yielding a true two-layer network whose hidden nonlinearity is derived, not designed.

We prove that the sigmoid's recurrence across both layers is a consequence of the exponential family structure of Bernoulli random variables, and independently derive the ReLU activation as the MAP estimator under sparse non-negative priors — establishing that the two dominant activations in deep learning answer complementary probabilistic questions: *"how probable?"* (sigmoid) and *"how much?"* (ReLU). We further prove that the Swish activation is the Bayesian expected value of relevant signal under the same sparse gating structure — the posterior mean counterpart to ReLU's posterior mode — establishing that the MAP-to-Bayes duality of classical statistics manifests as the ReLU-to-Swish transition, and that GELU is the Gaussian (probit) approximation of Swish. We show that WAND and Block-Max WAND algorithms from information retrieval constitute exact neural pruning methods with formal safety guarantees — a property enabled by sigmoid's boundedness and unattainable with ReLU's unboundedness, both properties now understood as consequences of the respective probabilistic questions.

Furthermore, relaxing the assumption of uniform signal reliability naturally extends the derived structure to the attention mechanism, identifying it as a form of context-dependent Logarithmic Opinion Pooling — mathematically equivalent to Hinton's Product of Experts (PoE, 2002) — and providing a probabilistic justification for why attention computes a weighted sum. We establish that the derived inference unit is the atomic building block of arbitrarily deep networks: depth arises from iterated marginalization over latent variables, where each layer constructs the evidence required by the next.

Our results reverse the conventional direction of explanation in neural network theory: rather than building neural networks and analyzing them probabilistically, we begin with probability and arrive at neurons, activations, attention, and depth. The correspondence between activation functions and probabilistic questions reframes architecture design as question sequencing — choosing the order of probabilistic questions posed to the data — and enables a new form of interpretability in which the activation function of each layer identifies the type of inference it performs. This provides an existence proof that neural architectures can arise as theorems of probabilistic reasoning, with immediate implications for network interpretability, efficient inference, and the theoretical foundations of neural computation.

1. Introduction

1.1 Background and Motivation

The relationship between probabilistic inference and neural computation has been a subject of sustained inquiry across multiple disciplines. The standard direction of investigation proceeds from neural networks to probabilistic interpretation: one constructs a neural architecture and subsequently asks what probabilistic model it corresponds to. Bayesian neural networks (Neal, 1996), variational inference methods (Blundell et al., 2015), and probabilistic deep learning (Gal & Ghahramani, 2016) all follow this direction.

In this paper, we demonstrate that the reverse direction is also productive — and yields a concrete, fully traceable result. We begin with a purely probabilistic question in information retrieval and show that the answer, when derived analytically, produces the computational structure of a feedforward neural network.

1.2 The Probabilistic Relevance Gap

Robertson (1977) introduced the Probability Ranking Principle (PRP), establishing that optimal document retrieval is achieved by ranking documents in decreasing order of their probability of relevance. Robertson and Zaragoza (2009) subsequently derived the BM25 scoring function from a probabilistic model of term occurrence, titling their foundational work *The Probabilistic Relevance Framework*.

Yet BM25 scores are not probabilities. The framework begins in probability theory and ends in unbounded real numbers. For nearly five decades, this gap persisted — acknowledged in standard textbooks (Manning et al., 2008; Croft et al., 2010), worked around in practice, never formally closed.

Problem 1.2.1 (The Probabilistic Relevance Gap). BM25 scores $s \in [0, +\infty)$ lack probabilistic interpretation, preventing principled combination with other signals and violating the premise of the Probability Ranking Principle.

In our companion paper (Jeong, 2026), we introduced Bayesian BM25 to close this gap, transforming BM25 scores into calibrated probability estimates through Bayesian inference with a sigmoid likelihood model and progressive hyperparameter estimation.

1.3 The Present Contribution

This paper takes the next step. We show that when multiple calibrated probability signals are combined through principled Bayesian reasoning, the resulting computational structure is not merely *analogous* to a neural network — it *is* one. Specifically, we prove:

1. **Analytical derivation of neural structure** (Section 5): The end-to-end computation for multi-signal probabilistic relevance estimation has the structure of a feedforward neural network with logit hidden nonlinearity and sigmoid output activation. When all signals share sigmoid calibration, the hidden layer collapses via the logit-sigmoid identity, yielding logistic regression (the Bayesian-logistic correspondence). When signals have heterogeneous calibrations — the practically important case of hybrid search — the logit performs a genuine nonlinear transformation, producing a true two-layer network.

2. **Inevitability of activation functions** (Section 6): The sigmoid's recurrence is a consequence of the Bernoulli exponential family. ReLU is independently derived as the MAP estimator under sparse non-negative priors. Swish is derived as the Bayesian expected value of relevant signal under the same sparse gating structure — the posterior mean counterpart to ReLU's posterior mode — establishing that the MAP-to-Bayes duality of classical statistics manifests as the ReLU-to-Swish transition in neural activations. GELU is proven to be the Gaussian (probit) approximation of Swish. The activations answer a hierarchy of probabilistic questions — "how probable?" (sigmoid), "how much?" (ReLU), and "what is the expected relevant amount?" (Swish) — explaining why hidden layers use ReLU or Swish and output layers use sigmoid.

3. **Exact neural pruning** (Section 7): WAND and Block-Max WAND algorithms from information retrieval constitute provably exact pruning methods for this neural structure, with formal safety guarantees that are unattainable with standard unbounded activations.

4. **From feedforward to attention** (Section 8): Relaxing the uniform reliability assumption in the derived network — allowing weights to depend on query-signal interaction — yields the attention mechanism as Logarithmic Opinion Pooling (equivalently, Hinton's Product of Experts) with context-dependent reliability.

5. **Depth, question sequencing, and interpretability** (Section 9): Deep networks are chains of recursive Bayesian inference over latent variables, where each layer constructs the evidence required by subsequent layers. The correspondence between activation functions and probabilistic questions reframes architecture design as question sequencing and enables a reverse interpretability method identifying the type of inference each layer performs.

6. **Reversal of explanatory direction** (Section 10): The derivation establishes that the neural structure is a *consequence* of probabilistic inference rather than a *design decision*, with implications for interpretability and the theoretical foundations of neural computation.

1.4 Notation

Throughout this paper, we use the following notation:

Symbol	Definition
$\sigma(x)$	Sigmoid function: $\frac{1}{1+\exp(-x)}$
$\text{logit}(p)$	Log-odds function: $\log \frac{p}{1-p}$
R	Binary relevance variable, $R \in \{0, 1\}$
s_i	Raw score from the i-th scoring signal
P_i	Calibrated probability from the i-th signal
$\bar{\ell}$	Log-odds mean of calibrated probabilities
n	Number of scoring signals
α_i, β_i	Sigmoid parameters for signal i
α	Confidence scaling exponent ($\alpha = 0.5$ yields $\sqrt{n}$ law)
λ	Exponential prior rate parameter (Section 6.5)
τ	Gaussian noise standard deviation (Section 6.5)

2. Mathematical Preliminaries

2.1 The Sigmoid and Logit Functions

Definition 2.1.1 (Sigmoid Function). The sigmoid function $\sigma : \mathbb{R} \to (0, 1)$ is defined as:

$$\sigma(x) = \frac{1}{1 + \exp(-x)} \tag{1}$$

Definition 2.1.2 (Logit Function). The logit function $\text{logit} : (0, 1) \to \mathbb{R}$ is the inverse of the sigmoid:

$$\text{logit}(p) = \log \frac{p}{1 - p} \tag{2}$$

Lemma 2.1.3 (Sigmoid Properties). The sigmoid function satisfies:

(i) *Symmetry*: $\sigma(-x) = 1 - \sigma(x)$

(ii) *Self-referential derivative*: $\sigma'(x) = \sigma(x)(1 - \sigma(x))$

(iii) *Bounds*: $\lim_{x \to \infty} \sigma(x) = 1$ and $\lim_{x \to -\infty} \sigma(x) = 0$

(iv) *Strict monotonicity*: $\sigma'(x) > 0$ for all $x \in \mathbb{R}$

Proof. Property (i):

$$\sigma(-x) = \frac{1}{1 + e^x} = \frac{e^{-x}}{e^{-x} + 1} = 1 - \frac{1}{1 + e^{-x}} = 1 - \sigma(x) \tag{3}$$

Property (ii): Let $L = \sigma(x) = (1 + e^{-x})^{-1}$. Then:

$$\sigma'(x) = e^{-x}(1 + e^{-x})^{-2} = \frac{1}{1 + e^{-x}} \cdot \frac{e^{-x}}{1 + e^{-x}} = \sigma(x)(1 - \sigma(x)) \tag{4}$$

Property (iv) follows from (ii), since $\sigma(x) \in (0, 1)$ implies $\sigma(x)(1 - \sigma(x)) > 0$. $\square$

Lemma 2.1.4 (Logit-Sigmoid Duality). For any $p \in (0, 1)$:

$$\sigma(\mathrm{logit}(p)) = p \quad \text{and} \quad \mathrm{logit}(\sigma(x)) = x \tag{5}$$

Proof. Direct computation:

$$\sigma(\mathrm{logit}(p)) = \sigma\left(\log \frac{p}{1 - p}\right) = \frac{1}{1 + \exp\left(-\log \frac{p}{1-p}\right)} = \frac{1}{1 + \frac{1-p}{p}} = p \quad \square \tag{6}$$

2.2 The Exponential Family and Canonical Links

Definition 2.2.1 (Exponential Family). A probability distribution belongs to the exponential family if its density can be written as:

$$f(x \mid \theta) = h(x) \exp(\eta(\theta) \cdot T(x) - A(\theta)) \tag{7}$$

where $\eta(\theta)$ is the natural parameter, $T(x)$ is the sufficient statistic, and $A(\theta)$ is the log-partition function.

Proposition 2.2.2 (Bernoulli Canonical Link). The Bernoulli distribution $\mathrm{Ber}(p)$ belongs to the exponential family with natural parameter $\eta = \log \frac{p}{1-p} = \mathrm{logit}(p)$. The canonical link function mapping the mean parameter p to the natural parameter η is the logit function, and its inverse — the mean function mapping η to p — is the sigmoid.

Proof. The Bernoulli PMF is:

$$P(x \mid p) = p^x (1 - p)^{1-x} = (1 - p) \exp\left(x \log \frac{p}{1 - p}\right) \tag{8}$$

This is in exponential family form with $\eta = \mathrm{logit}(p)$, $T(x) = x$, $h(x) = 1$, and $A(\eta) = \log(1 + e^\eta)$. The mean is:

$$\mathbb{E}[X] = A'(\eta) = \frac{e^\eta}{1 + e^\eta} = \sigma(\eta) \quad \square \tag{9}$$

2.3 Bayesian Inference for Binary Relevance

Definition 2.3.1 (Bayesian Posterior for Relevance). For a binary relevance variable $R \in \{0, 1\}$ and observed score s, the posterior probability of relevance is given by Bayes' theorem:

$$P(R = 1 \mid s) = \frac{P(s \mid R = 1) \cdot P(R = 1)}{P(s \mid R = 1) \cdot P(R = 1) + P(s \mid R = 0) \cdot P(R = 0)} \tag{10}$$

Definition 2.3.2 (Sigmoid Likelihood Model). We model the likelihood ratio of observing score s as:

$$\Lambda(s) = \frac{P(s \mid R = 1)}{P(s \mid R = 0)} = e^{\alpha(s - \beta)} \tag{11}$$

The **normalized likelihood ratio** is the sigmoid transformation:

$$L(s) = \frac{\Lambda(s)}{1 + \Lambda(s)} = \sigma(\alpha(s - \beta)) \tag{12}$$

Remark 2.3.2a (Measure-Theoretic Foundation). The log-likelihood ratio $\log \Lambda(s) = \alpha(s - \beta)$ is a linear function of s, which is the canonical form for exponential family models. This formulation does not require $L(s)$ to be a probability density function — it is a Radon-Nikodym derivative ratio that quantifies the relative evidence for $R = 1$ versus $R = 0$. The sigmoid $L(s) = \sigma(\alpha(s - \beta))$ maps this ratio to $[0, 1]$ as a normalized belief, not as a density. See Jeong (2026, Section 4.1) for the complete derivation.

Definition 2.3.3 (Symmetric Likelihood Assumption). The log-linear likelihood ratio model implies:

$$P(s \mid R = 0) = 1 - P(s \mid R = 1) = \sigma(-\alpha(s - \beta)) \tag{13}$$

This follows from the algebraic identity $L(s) + (1 - L(s)) = 1$ and is equivalent to the assumption that the log-likelihood ratio is linear in s — a standard assumption in logistic regression and Platt scaling (Platt, 1999).

Theorem 2.3.4 (Sigmoid Posterior). Under the symmetric likelihood assumption (Definition 2.3.3) with uniform prior $P(R = 1) = P(R = 0) = 0.5$, the posterior probability of relevance reduces to:

$$P(R = 1 \mid s) = \sigma(\alpha(s - \beta)) \tag{14}$$

Proof. Substituting into Bayes' theorem:

$$P(R = 1 \mid s) = \frac{\sigma(\alpha(s - \beta)) \cdot 0.5}{\sigma(\alpha(s - \beta)) \cdot 0.5 + \sigma(-\alpha(s - \beta)) \cdot 0.5} \tag{15}$$

By the symmetry property $\sigma(-x) = 1 - \sigma(x)$:

$$= \frac{\sigma(\alpha(s - \beta))}{\sigma(\alpha(s - \beta)) + 1 - \sigma(\alpha(s - \beta))} = \sigma(\alpha(s - \beta)) \quad \square \tag{16}$$

Remark 2.3.5 The sigmoid was not chosen as a convenient functional form — it was *derived* as the unique posterior under the stated likelihood model. This derivation, established in Jeong (2026), is the foundation upon which the present paper builds.

---

3. The Conjunction Shrinkage Problem

3.1 Naive Probabilistic Conjunction

Given n independent calibrated relevance signals $P_1, P_2, \ldots, P_n$, the standard probabilistic conjunction under independence is:

Definition 3.1.1 (Product Rule Conjunction).

$$P_{\text{AND}} = \prod_{i=1}^{n} P_i \tag{17}$$

This formula, presented in Jeong (2026, Section 5.1), is theoretically correct under the stated independence assumptions. However, it suffers from a fundamental deficiency when applied to evidence accumulation.

3.2 The Shrinkage Theorem

Theorem 3.2.1 (Conjunction Shrinkage). For n independent signals each reporting probability $p \in (0, 1)$:

$$\prod_{i=1}^{n} p = p^n \xrightarrow{n \to \infty} 0 \tag{18}$$

Proof. Since $0 < p < 1$, we have $\log p < 0$, so $n \log p \to -\infty$ as $n \to \infty$, and $\exp(n \log p) \to 0$. $\square$

Corollary 3.2.2 (Monotone Shrinkage). The conjunction probability is strictly decreasing in n:

$$\prod_{i=1}^{n+1} p < \prod_{i=1}^{n} p \tag{19}$$

for all $p \in (0, 1)$ and $n \geq 1$.

3.3 The Semantic Mismatch

Problem 3.3.1 (Evidence Accumulation vs. Joint Satisfaction). The product rule answers the question "what is the probability that all conditions are simultaneously satisfied?" However, the question a search system poses is "how confident should we be given that multiple signals concur?" These are semantically distinct questions.

When n independent signals all report high relevance, the product rule yields a probability that decreases with n. This violates the fundamental intuition of evidence accumulation: *agreement among independent sources should increase confidence, not decrease it.*

3.4 Information-Theoretic Analysis

Proposition 3.4.1 (Information Loss in Product Conjunction). The product rule discards the mutual agreement information among signals. Specifically, for n i.i.d. signals each reporting probability p, the product p^n depends only on p and n, and is invariant to whether the signals agree or disagree — they are treated as independent filters rather than corroborating testimony.

Proof. The product $\prod_{i=1}^{n} P_i$ is symmetric in the P_i and contains no interaction terms. The agreement structure — whether all signals report similar values or diverse values — is not represented in the product. $\square$

4. Log-Odds Conjunction: A Framework for Evidence Accumulation

We now present a conjunction framework that resolves the shrinkage problem while preserving probabilistic soundness.

4.1 Log-Odds Mean Aggregation

The geometric mean $\bar{P} = (\prod P_i)^{1/n}$ resolves the shrinkage problem in probability space by preserving scale neutrality: $\bar{P} = p$ when $P_i = p$ for all i (see Section 3). However, the subsequent transformation to log-odds introduces a nonlinear residual (see Remark 4.1.3 below). We therefore work directly in the log-odds domain, which yields an exact linear aggregation.

Definition 4.1.1 (Log-Odds Mean). The log-odds mean of n calibrated probabilities $P_1, \ldots, P_n$ is the arithmetic mean of their logits:

$$\bar{\ell} = \frac{1}{n} \sum_{i=1}^{n} \mathrm{logit}(P_i) = \frac{1}{n} \sum_{i=1}^{n} \log \frac{P_i}{1 - P_i} \tag{20}$$

Theorem 4.1.2 (Scale Neutrality in Log-Odds). If $P_i = p$ for all i, then $\bar{\ell} = \mathrm{logit}(p)$ and $\sigma(\bar{\ell}) = p$, regardless of n.

Proof. $\bar{\ell} = \frac{1}{n} \sum_{i=1}^{n} \mathrm{logit}(p) = \mathrm{logit}(p)$. By Lemma 2.1.4, $\sigma(\mathrm{logit}(p)) = p$. $\square$

Theorem 4.1.2a (Equivalence to Normalized Log-OP). The log-odds mean is the exact normalized form of the Logarithmic Opinion Pool (Log-OP), also known as the Product of Experts (PoE, Hinton, 2002). Given uniform weights $w_i = 1/n$:

$$P_{\text{Log-OP}} = \frac{\prod_i P_i^{1/n}}{\prod_i P_i^{1/n} + \prod_i (1 - P_i)^{1/n}} \tag{21}$$

Taking the logit of both sides:

$$\mathrm{logit}(P_{\text{Log-OP}}) = \sum_{i=1}^{n} \frac{1}{n} \log \frac{P_i}{1 - P_i} = \frac{1}{n} \sum_{i=1}^{n} \mathrm{logit}(P_i) = \bar{\ell} \tag{22}$$

The log-odds mean is therefore the logit of the normalized Log-OP — the exact representation of Product-of-Experts aggregation in the natural parameter space of Bernoulli random variables. $\square$

Remark 4.1.3 (Why Not Geometric Mean Followed by Logit). The geometric mean $\bar{P} = (\prod P_i)^{1/n}$ operates in log-probability space: $\log \bar{P} = \frac{1}{n} \sum \log P_i$. Converting this to log-odds requires the nonlinear transformation $\mathrm{logit}(\exp(S)) = S - \log(1 - e^S)$ where $S = \log \bar{P}$. The residual term $-\log(1 - e^S)$ vanishes only when $\bar{P} \ll 1$, introducing approximation error for relevant documents ($\bar{P} \gg 0$). Working directly with logits avoids this issue entirely: the Log-OP is *exactly* linear in the logit domain, with no approximation at any operating point.

Remark 4.1.4 The log-odds mean neutralizes the dependence on signal count that causes conjunction shrinkage (Theorem 3.2.1). It preserves the "average log-odds evidence" of the constituent signals without penalizing for their number, and is the unique aggregation that is both scale-neutral and exactly linear in the natural parameter of the Bernoulli family.

4.2 Multiplicative Confidence Scaling

Definition 4.2.1 (Log-Odds Conjunction). Given the log-odds mean $\bar{\ell}$, the confidence-scaled log-odds conjunction is:

$$\ell_{\text{adjusted}} = \bar{\ell} \cdot n^\alpha = \left(\frac{1}{n} \sum_{i=1}^{n} \mathrm{logit}(P_i) \right) \cdot n^\alpha \tag{23}$$

where $\alpha \geq 0$ is a scaling constant. Equivalently:

$$\ell_{\text{adjusted}} = \frac{1}{n^{1-\alpha}} \sum_{i=1}^{n} \mathrm{logit}(P_i) \tag{24}$$

Theorem 4.2.2 (Sign Preservation). The multiplicative scaling preserves the sign of the log-odds mean:

$$\text{sgn}(\ell_{\text{adjusted}}) = \text{sgn}(\bar{\ell}) \tag{25}$$

Proof. Since $n^\alpha > 0$ for all $n \geq 1$ and $\alpha \geq 0$, multiplying $\bar{\ell}$ by n^α preserves its sign. $\square$

Corollary 4.2.3 (Irrelevance Cannot Be Inverted). If $\text{logit}(P_i) < 0$ for all i (all signals report irrelevance), then $\bar{\ell} < 0$ and therefore $\ell_{\text{adjusted}} < 0$, which implies $P_{\text{final}} = \sigma(\ell_{\text{adjusted}}) < 0.5$. No amount of agreement among irrelevant signals can produce a relevance judgment. This is a structural guarantee absent from additive bias formulations.

Remark 4.2.4 (Why Multiplicative, Not Additive). An additive conjunction bonus $\bar{\ell} + \alpha \log n$ would add a positive constant to the log-odds regardless of sign. For sufficiently large n, this constant can dominate a negative $\bar{\ell}$, causing $P_{\text{final}} > 0.5$ even when all signals agree on irrelevance — a catastrophic violation of evidence accumulation semantics. The multiplicative form $\bar{\ell} \cdot n^\alpha$ amplifies the *magnitude* of the log-odds while preserving its *direction*, correctly modeling the intuition that "agreement among n sources strengthens the conclusion in whatever direction the evidence points."

4.3 Return to Probability Space

Definition 4.3.1 (Final Posterior). The final combined probability is obtained by applying the inverse logit (sigmoid) to the scaled log-odds:

$$P_{\text{final}} = \sigma(\ell_{\text{adjusted}}) = \sigma\left(\frac{1}{n^{1-\alpha}} \sum_{i=1}^{n} \text{logit}(P_i) \right) \tag{26}$$

Proposition 4.3.2 (Identity for Single Signals). When $n = 1$, the scaling factor is $n^\alpha = 1$:

$$\ell_{\text{adjusted}} = \text{logit}(P_1) \cdot 1 = \text{logit}(P_1) \implies P_{\text{final}} = \sigma(\text{logit}(P_1)) = P_1 \tag{27}$$

The transformation is transparent for single signals.

Proof. Follows from Lemma 2.1.4 (logit-sigmoid duality). $\square$

4.4 The $\sqrt{n}$ Scaling Law

Theorem 4.4.1 (Statistical Justification for $\alpha = 0.5$). Setting $\alpha = 0.5$ yields a weight of $w_i = 1/\sqrt{n}$ per signal, embedding the classical $\sqrt{n}$ confidence scaling law directly in the log-odds domain:

$$\ell_{\text{adjusted}} = \frac{1}{\sqrt{n}} \sum_{i=1}^{n} \text{logit}(P_i) \tag{28}$$

Justification. In classical statistics, when combining n independent measurements of the same quantity, the standard error of the mean decreases proportionally to $1/\sqrt{n}$, and the corresponding test statistic (z-score, t-statistic) increases proportionally to $\sqrt{n}$. The log-odds of each signal $\text{logit}(P_i)$ is a measure of evidence strength. The total evidence from n independent signals is $\sum \text{logit}(P_i)$, but the *averaged* evidence is $\bar{\ell} = \frac{1}{n} \sum \text{logit}(P_i)$. The $\sqrt{n}$ scaling bridges these two extremes: it weights each signal at $1/\sqrt{n}$, producing a total that grows as $\sqrt{n} \cdot \bar{\ell}$ — exactly the rate at which confidence should grow under independent observations. $\square$

Proposition 4.4.2 (Confidence Scaling Magnitudes). For $\alpha = 0.5$, the effective weight per signal is:

n (signals)	Weight $w_i = 1/\sqrt{n}$	Total weight $\sqrt{n}$
1	1.000	1.00
2	0.707	1.41
3	0.577	1.73
5	0.447	2.24
10	0.316	3.16

4.5 Numerical Behavior

Theorem 4.5.1 (Behavioral Properties). The log-odds conjunction satisfies the following properties:

(i) *Agreement amplification*: If $P_i > 0.5$ for all i, then $P_{\text{final}} > \sigma(\bar{\ell})$ for $n \geq 2$.

(ii) *Disagreement moderation*: If signals disagree symmetrically (logits sum to near zero), then $P_{\text{final}} \approx 0.5$ (exact neutrality when logits cancel).

(iii) *Irrelevance preservation*: If $P_i < 0.5$ for all i, then $P_{\text{final}} < 0.5$ for all n.

(iv) *Relevance preservation*: If $P_i > 0.5$ for all i, then $P_{\text{final}} > 0.5$ for all n.

Proof of (iii). If $P_i < 0.5$ for all i, then $\text{logit}(P_i) < 0$ for all i. Therefore $\bar{\ell} = \frac{1}{n} \sum \text{logit}(P_i) < 0$. Since $n^\alpha > 0$, the product $\ell_{\text{adjusted}} = \bar{\ell} \cdot n^\alpha < 0$, and $P_{\text{final}} = \sigma(\ell_{\text{adjusted}}) < \sigma(0) = 0.5$. $\square$

Proof of (i). If $\bar{\ell} > 0$ (all $P_i > 0.5$), then $\ell_{\text{adjusted}} = \bar{\ell} \cdot n^\alpha > \bar{\ell}$ for $n \geq 2$ and $\alpha > 0$. By strict monotonicity of σ, $P_{\text{final}} = \sigma(\ell_{\text{adjusted}}) > \sigma(\bar{\ell})$. $\square$

Remark 4.5.2 (Structural Guarantee). Properties (iii) and (iv) are structural consequences of the multiplicative formulation (Theorem 4.2.2) — they hold for *all n*, *all $\alpha \geq 0$*, and *all* configurations of P_i. They cannot be violated by any parameter choice. This contrasts with additive bias formulations where sufficiently large n can invert the sign of the evidence.

The following table compares the product rule and log-odds conjunction for two signals ($n = 2$, $\alpha = 0.5$):

P_{text}	P_{vec}	Product	Log-Odds Conjunction	Interpretation
0.9	0.9	0.81	0.96	Strong agreement amplified
0.7	0.7	0.49	0.77	Moderate agreement preserved
0.7	0.3	0.21	0.50	Exact neutrality (logits cancel)
0.3	0.3	0.09	0.23	Irrelevance preserved

5. The Emergence of Neural Network Structure

We now arrive at the central result of this paper.

5.1 The General Computational Pipeline

We trace the full computation for estimating the probability that a document d is relevant given n calibrated probability signals in a hybrid search query.

Stage 1 — Calibration. Each scoring signal produces a calibrated probability $P_i \in (0, 1)$ representing the probability that d is relevant according to signal i. The calibration method may differ across signals:

- **Sigmoid calibration** (BM25 and similar raw scores): $P_i = \sigma(\alpha_i(s_i - \beta_i))$, where α_i, β_i are estimated from the score distribution (Jeong, 2026, Section 4).

- **Linear calibration** (cosine similarity): $P_i = (1 + s_i)/2$, mapping $s_i \in [-1, 1]$ to $[0, 1]$ (Jeong, 2026, Definition 7.1.2).

- **External calibration** (pre-trained neural rankers, cross-encoders): P_i is the model's output probability, calibrated by its own training procedure.

The critical property is that after Stage 1, all signals share a common representation: probabilities in $(0, 1)$.

Stage 2 — Log-Odds Aggregation. The calibrated probabilities are mapped to log-odds and aggregated as a weighted sum (normalized Log-OP, Theorem 4.1.2a):

$$\bar{\ell} = \frac{1}{n} \sum_{i=1}^{n} \text{logit}(P_i) = \frac{1}{n} \sum_{i=1}^{n} \log \frac{P_i}{1 - P_i} \tag{29}$$

Stage 3 — Confidence Scaling and Final Posterior. The mean log-odds is scaled by n^α and passed through a sigmoid to produce the final probability:

$$P_{\text{final}} = \sigma\left(\bar{\ell} \cdot n^\alpha\right) = \sigma\left(\frac{1}{n^{1-\alpha}} \sum_{i=1}^{n} \text{logit}(P_i)\right) \tag{30}$$

5.2 Identification of Neural Structure

Theorem 5.2.1 (Neural Network Structure). The computation described in Section 5.1 has the structure of a two-layer feedforward network: an input-to-probability calibration layer, a hidden transformation via the logit function, a weighted sum in log-odds space, and an output sigmoid activation. This structure is exact at all operating points.

Proof. We identify the components:

(i) **Input layer**: Scoring signals producing probabilities $P_1, P_2, \ldots, P_n$.

(ii) **Hidden nonlinearity**: Each probability P_i is mapped through the logit function to produce hidden-layer inputs:

$$x_i = \text{logit}(P_i) = \log \frac{P_i}{1 - P_i} \tag{31}$$

The logit is a monotonic nonlinear transformation that maps $(0, 1)$ to $\mathbb{R}$, converting probabilities to log-odds — the natural parameter of the Bernoulli exponential family (Proposition 2.2.2).

(iii) **Linear aggregation**: The hidden-layer outputs are combined as a weighted sum:

$$\ell_{\text{adjusted}} = \sum_{i=1}^{n} w_i \cdot x_i = \frac{1}{n^{1-\alpha}} \sum_{i=1}^{n} \text{logit}(P_i) \tag{32}$$

where $w_i = 1/n^{1-\alpha}$.

(iv) **Output activation**: The result passes through a sigmoid σ to produce $P_{\text{final}} \in (0, 1)$.

The full computation is:

$$P_{\text{final}} = \sigma\left(\frac{1}{n^{1-\alpha}} \sum_{i=1}^{n} \text{logit}(P_i) \right) \tag{33}$$

This is a feedforward network with logit nonlinearity in the hidden layer, linear aggregation, and sigmoid output activation. □

The network's hidden nonlinearity is the logit function — the canonical link of the Bernoulli exponential family. Unlike ReLU or tanh, this nonlinearity was not selected from a design space; it was derived as the unique function that maps probabilities to the natural parameter space where evidence combination is additive (Theorem 4.1.2a).

5.2a The Sigmoid-Calibrated Special Case

When all signals are sigmoid-calibrated from raw scores — the case where Stage 1 uses $P_i = \sigma(\alpha_i s_i + \beta_i')$ for all i — a remarkable algebraic simplification occurs.

Theorem 5.2.1a (Logistic Regression Equivalence). If all calibrated probabilities are produced by sigmoid functions of their respective raw scores, i.e., $P_i = \sigma(\alpha_i s_i + \beta_i')$ where $\beta_i' = -\alpha_i \beta_i$, then the hidden nonlinearity collapses via the logit-sigmoid identity, and the full pipeline reduces to a single-neuron logistic regression:

$$P_{\text{final}} = \sigma\left(\sum_{i=1}^{n} w_i' \cdot s_i + b \right) \tag{34}$$

where $w_i' = \alpha_i/n^{1-\alpha}$ are the effective weights and $b = \sum_i \beta_i'/n^{1-\alpha}$ is the effective bias.

Proof. When $P_i = \sigma(\alpha_i s_i + \beta_i')$, the logit-sigmoid duality (Lemma 2.1.4) yields:

$$\text{logit}(P_i) = \text{logit}(\sigma(\alpha_i s_i + \beta_i')) = \alpha_i s_i + \beta_i' \tag{35}$$

The hidden nonlinearity — the logit — acts as the identity on sigmoid-calibrated inputs. Substituting into the general formula:

$$P_{\text{final}} = \sigma\left(\frac{1}{n^{1-\alpha}} \sum_{i=1}^{n} (\alpha_i s_i + \beta_i') \right) = \sigma\left(\sum_{i=1}^{n} w_i' \cdot s_i + b \right) \quad \square \tag{36}$$

Remark 5.2.2 (Bayesian Inference as Logistic Regression). Theorem 5.2.1a is not a weakness — it is a classical result seen from a new angle. The equivalence between Bayesian posterior computation under sigmoid likelihood and logistic regression has been known since at least Cox (1958) and Berkson (1944). Our derivation arrives at this correspondence from a different starting point — evidence accumulation for multi-signal retrieval — confirming that the logistic regression form is not an approximation to Bayesian inference but *is* Bayesian inference when signals share the same exponential family calibration.

Remark 5.2.3 (Why the Two-Layer Structure is Genuine for Heterogeneous Signals). The logit-sigmoid cancellation of Theorem 5.2.1a occurs *only* when $P_i = \sigma(\text{linear}(s_i))$ — that is, when every signal's calibration belongs to the same parametric family as the output activation. In practice, hybrid search combines signals with *different* calibrations:

- **BM25**: $P_1 = \sigma(\alpha_1 s_1 + \beta_1')$ — sigmoid calibration (logit recovers linear pre-activation)

- **Vector similarity**: $P_2 = (1 + \cos\theta)/2$ — linear calibration (Jeong, 2026, Definition 7.1.2)

For the linearly-calibrated vector signal:

$$\text{logit}(P_2) = \text{logit}\left(\frac{1 + \cos\theta}{2}\right) = \log\frac{1 + \cos\theta}{1 - \cos\theta} \tag{37}$$

This is a *nonlinear* function of $\cos\theta$: the logit of a linear map is not linear. The hidden layer therefore performs a genuine nonlinear transformation on the vector signal, even though it acts as the identity on the BM25 signal. The network has:

- A **linear pathway** for sigmoid-calibrated signals (logit-sigmoid cancellation)

- A **nonlinear pathway** for differently-calibrated signals (logit as genuine nonlinearity)

This mixed structure — identity skip connections for some inputs, nonlinear transformation for others — arises naturally from the heterogeneity of the calibration methods, not from architectural design.

Remark 5.2.4 (The Logit-Sigmoid Cancellation as Calibration Transparency). The cancellation in Theorem 5.2.1a has a probabilistic interpretation: when a signal's calibration is already in the exponential family's natural parametrization (sigmoid), the conversion to log-odds adds no information — the signal was already in the correct coordinate system. Different calibrations (linear, probit, learned) require a genuine coordinate change via the logit, which is the source of the hidden-layer nonlinearity. The logit function thus acts as a **calibration normalizer**: it is transparent to correctly parametrized signals and applies a corrective nonlinearity to all others.

5.3 Parameter Correspondence

Theorem 5.3.1 (Explicit Parameter Mapping). The correspondence between the probabilistic derivation and neural network parameters is:

Probabilistic Component	Neural Network Component
Calibrated probabilities P_i	First-layer outputs
$\text{logit}(P_i)$	Hidden-layer activations (logit nonlinearity)
Confidence-scaled weights $w_i = 1/n^{1-\alpha}$	Hidden-layer weights
Aggregated bias $b = \sum_i w_i \cdot \text{logit}(P_i)\vert_{P_i \text{ const.}}$	Output neuron bias contribution
Final sigmoid $\sigma(\cdot)$	Output-layer sigmoid activation
P_{final}	Network output

In the sigmoid-calibrated special case (Theorem 5.2.1a), the correspondence simplifies further: the effective weights are $w_i' = \alpha_i/n^{1-\alpha}$ and the effective bias is $b = \sum_i \beta_i'/n^{1-\alpha}$, reducing to a single logistic regression neuron.

Remark 5.3.2 In the uniform case, the weights $w_i = 1/n^{1-\alpha}$ are determined by the number of signals and the confidence scaling parameter. In the weighted case, where signals have different reliabilities, the weights w_i become learnable parameters, completing the correspondence to a fully parameterized neural network.

5.4 Directionality of Derivation

Remark 5.4.1 (Derivation vs. Design). The neural structure identified in Theorem 5.2.1 was not designed to resemble a neural network. At no point in the derivation — from Bayes' theorem (Section 2.3) through conjunction shrinkage resolution (Section 4) to the final posterior computation — was any architectural decision made with neural computation in mind. The structure emerged as a consequence of the mathematics.

The derivation reveals two complementary results. First, when all signals are sigmoid-calibrated, the computation reduces to logistic regression (Theorem 5.2.1a) — confirming from a new direction that Bayesian posterior computation under the Bernoulli exponential family *is* logistic regression. Second, when signals arrive from heterogeneous calibrations — the practically important case of hybrid search — the computation is a genuine two-layer network with the logit as hidden nonlinearity (Remark 5.2.3). In both cases, the structure was derived, not designed.

---

6. The Inevitability of Activation Functions

6.1 Dual Appearance

Observation 6.1.1 The sigmoid function appears twice in the derivation:
(i) In Stage 1, as the Bayesian posterior for individual signal calibration (Theorem 2.3.4).
(ii) In Stage 3, as the inverse logit mapping adjusted log-odds back to probability space (Definition 4.3.1).

Neither appearance is an architectural choice. Both are mathematical necessities.

6.2 Characterization Theorem

Theorem 6.2.1 (Uniqueness of the Sigmoid). The logistic sigmoid is the unique function satisfying the following system of constraints simultaneously:

(C1) $\sigma : \mathbb{R} \to (0, 1)$ — maps real-valued inputs to valid probabilities.

(C2) $\sigma(x) = [\text{logit}]^{-1}(x)$ — is the canonical inverse link for the Bernoulli exponential family.

(C3) $\sigma'(x) = \sigma(x)(1 - \sigma(x))$ — has a self-referential derivative expressible in terms of its own output.

(C4) $\sigma(-x) = 1 - \sigma(x)$ — is symmetric with respect to positive and negative evidence.

(C5) σ arises as the maximum entropy distribution for binary outcomes under first-moment constraints.

Proof sketch. Constraint (C2) uniquely determines σ within the exponential family framework (Proposition 2.2.2). Constraints (C1), (C3), and (C4) follow as consequences (Lemma 2.1.3). Constraint (C5) follows from the fact that the Bernoulli distribution is the maximum entropy distribution over $\{0, 1\}$ given a specified mean, and the logistic function is its natural parametrization. $\square$

Corollary 6.2.2 (Inevitability). Any system that processes binary evidence — relevant or irrelevant, firing or quiescent, true or false — and respects the constraints (C1)–(C5) will arrive at the sigmoid function. The sigmoid's appearance in both Bayesian inference and neural computation is not a coincidence but a consequence of the shared mathematical structure of binary probabilistic reasoning.

6.3 Exclusion of Alternative Activation Functions

To clarify the strength of Theorem 6.2.1, we examine why commonly used activation functions fail to satisfy the constraint system (C1)–(C5).

Proposition 6.3.1 (Exclusion of $\tanh$). The hyperbolic tangent $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$ maps $\mathbb{R} \to (-1, 1)$, violating constraint (C1). The rescaled form $\frac{1}{2}(\tanh(x) + 1)$ maps to $(0, 1)$, but this identity holds because $\frac{1}{2}(\tanh(x) + 1) \equiv \sigma(2x)$. Any valid rescaling of $\tanh$ to the unit interval reduces identically to the sigmoid with parameter absorption. The $\tanh$ is not an independent alternative — it is the sigmoid in disguise.

Proof. The canonical link inverse requires $g^{-1}(\text{logit}(p)) = p$ for all $p \in (0, 1)$. For $g^{-1}(x) = \frac{1}{2}(\tanh(x) + 1)$:

$$g^{-1}(\text{logit}(p)) = \frac{1}{2}\left(\tanh\left(\log \frac{p}{1-p}\right) + 1\right) = p \tag{38}$$

This holds precisely because the rescaled $\tanh$ *is* $\sigma(2x)$, confirming rather than challenging the sigmoid's uniqueness. $\square$

Proposition 6.3.2 (Exclusion of Softplus). The softplus function $f(x) = \log(1 + e^x)$ maps $\mathbb{R} \to (0, +\infty)$, violating constraint (C1). Its output has no upper bound and therefore cannot represent a probability. Additionally, $f(-x) \neq 1 - f(x)$, violating the evidence symmetry constraint (C4).

Proposition 6.3.3 (Exclusion of ReLU). The ReLU function $f(x) = \max(0, x)$ maps $\mathbb{R} \to [0, +\infty)$, violating constraint (C1). It is not differentiable at $x = 0$, precluding a self-referential derivative (C3). It violates symmetry (C4) since $f(-x) = 0 \neq 1 - f(x)$ for $x > 0$. It does not arise from any exponential family canonical link (C2).

Proposition 6.3.4 (Exclusion of Probit). The probit function $\Phi(x) = \int_{-\infty}^{x} \frac{1}{\sqrt{2\pi}} e^{-t^2/2} \, dt$ satisfies constraints (C1) and (C4), and maps $\mathbb{R} \to (0, 1)$ with the symmetry $\Phi(-x) = 1 - \Phi(x)$. However, it does not arise as the canonical link inverse for the Bernoulli exponential family (C2), and its derivative $\Phi'(x) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2}$ cannot be expressed as a function of $\Phi(x)$ alone, violating (C3). The probit corresponds to Gaussian-distributed latent variables rather than the natural parametrization of Bernoulli outcomes.

Remark 6.3.5 (Summary of Exclusions). The following table summarizes the constraint violations:

Function	(C1) $\mathbb{R} \to (0,1)$	(C2) Canonical link	(C3) Self-ref. derivative	(C4) Symmetry	(C5) Max entropy
Sigmoid σ	✓	✓	✓	✓	✓
$\tanh$	✗	Reduces to σ	—	—	—
Softplus	✗	✗	✗	✗	✗
ReLU	✗	✗	✗	✗	✗
Probit Φ	✓	✗	✗	✓	✗

The sigmoid is the unique function satisfying all five constraints. Functions that satisfy a subset — most notably the probit, which satisfies (C1) and (C4) — fail on the exponential family constraints (C2), (C3), and (C5) that are essential for Bayesian inference over Bernoulli outcomes.

6.4 The Identity of Neuron and Posterior

Theorem 6.4.1 (Neuron-Posterior Identity). The Bayesian posterior for binary relevance under the sigmoid likelihood model and the output of a sigmoid neuron are the same mathematical object — the natural parameter-to-mean mapping in the Bernoulli exponential family.

Proof. By Proposition 2.2.2, the mean of a Bernoulli random variable as a function of its natural parameter η is $\sigma(\eta)$. The Bayesian posterior (Theorem 2.3.4) computes $\sigma(\alpha(s - \beta))$, which maps a score-derived natural parameter to a probability. A sigmoid neuron computes $\sigma(wx + b)$, which maps a weighted input to an activation in $(0, 1)$. Both are instances of $\sigma : \eta \mapsto p$ in the Bernoulli exponential family. $\square$

Remark 6.4.2 The neuron does not imitate the Bayesian posterior. The Bayesian posterior does not imitate the neuron. Both are the sigmoid because nothing else can satisfy the constraints of binary probabilistic reasoning.

Remark 6.4.3 (Extension to Multi-Class: Softmax). While this paper focuses on binary relevance, the framework extends naturally. For multi-class classification over K categories, the Bernoulli distribution generalizes to the Categorical distribution, which belongs to the exponential family with natural parameters $\eta_k = \log \frac{p_k}{p_K}$. The canonical inverse link mapping natural parameters to class probabilities is the softmax function $p_k = \frac{e^{\eta_k}}{\sum_j e^{\eta_j}}$. This confirms that the standard neural activation functions for both binary and multi-class outputs — sigmoid and softmax — are not engineering conveniences but manifestations of exponential family structure.

6.5 ReLU as MAP Estimation under Sparse Priors

Sections 6.1–6.4 established that the sigmoid is the inevitable answer to: *"What is the probability that the hypothesis is true?"* We now show that a different probabilistic question yields a different activation function with equal inevitability.

Observation 6.5.1 (A Different Question). Consider the question: *"How much of a latent feature h is present in the observation?"* This asks for a **quantity**, not a probability. The answer is non-negative (feature presence cannot be negative) and unbounded (there is no maximum amount of evidence).

Definition 6.5.2 (Sparse Feature Model). Let $h \geq 0$ be a latent variable representing the activation level of a feature, with:

(i) **Sparse prior** (Exponential distribution):

$$P(h) = \lambda e^{-\lambda h}, \quad h \geq 0 \tag{39}$$

(ii) **Gaussian likelihood**:

$$P(x \mid h) \propto \exp\left(-\frac{(x - wh)^2}{2\tau^2}\right) \tag{40}$$

The exponential prior encodes the assumption that most features are absent most of the time — the continuous analog of the term frequency distribution in information retrieval, where most terms are absent from most documents.

Theorem 6.5.3 (ReLU from MAP Estimation). The MAP estimate of h under the sparse feature model (Definition 6.5.2) is:

$$h^* = \max(0, \ z - \theta) \tag{41}$$

where $z = x/w$ is the normalized input and $\theta = \lambda\tau^2/w^2$ is a threshold determined by the prior strength and noise level. This is the ReLU activation function with bias $b = -\theta$.

Proof. The log-posterior is:

$$\mathcal{L}(h) = \log P(x \mid h) + \log P(h) = -\frac{(x - wh)^2}{2\tau^2} - \lambda h + \text{const} \tag{42}$$

Differentiating with respect to h:

$$\frac{\partial \mathcal{L}}{\partial h} = \frac{w(x - wh)}{\tau^2} - \lambda \tag{43}$$

Setting to zero yields the unconstrained optimum:

$$h_{\text{unc}} = \frac{wx - \lambda\tau^2}{w^2} \tag{44}$$

Applying the non-negativity constraint $h \geq 0$:

$$h^* = \max\left(0, \ \frac{wx - \lambda\tau^2}{w^2}\right) = \max(0, \ z - \theta) \quad \square \tag{45}$$

Theorem 6.5.4 (Characterization of ReLU). The ReLU form is the unique MAP estimator satisfying:

(Q1) **Non-negativity**: $h^* \geq 0$ (feature presence cannot be negative).

(Q2) **Sparsity**: $h^* = 0$ for a positive-measure set of inputs (most features are absent).

(Q3) **Linearity above threshold**: For sufficiently strong inputs, h^* grows linearly (evidence scales proportionally with signal strength).

(Q4) **Hard thresholding**: Below a critical input level, the output is exactly zero — not approximately zero, but structurally zero.

Proof sketch. Constraint (Q1) restricts the estimator to $[0, +\infty)$. Constraint (Q2) requires a threshold below which the output is identically zero. Constraint (Q3) requires linearity above the threshold. Together, (Q1)–(Q3) determine the form $\max(0, az + b)$ up to positive scaling. The exponential prior is the maximum entropy distribution on $[0, +\infty)$ given a specified mean (paralleling the role of the Bernoulli distribution as maximum entropy on $\{0, 1\}$), and its MAP solution uniquely satisfies (Q4). $\square$

Remark 6.5.5 (Why ReLU is Unbounded). The unboundedness of ReLU is not a defect — it is a direct consequence of the question being asked. *"How much?"* has no upper limit, just as *"How probable?"* has the natural upper limit of 1. The range of the activation function is determined by the semantics of the answer:

- Sigmoid answers "how probable?" → bounded in $(0, 1)$
- ReLU answers "how much?" → unbounded in $[0, +\infty)$

6.6 The Complementarity of Sigmoid, ReLU, and Swish

Proposition 6.6.1 (Three Questions, Three Activations). The sigmoid, ReLU, and Swish are not competing activation functions but answers to complementary probabilistic questions:

Property	Sigmoid	ReLU	Swish
Question	Is the hypothesis true?	How much evidence is present?	What is the expected relevant signal?
Output semantics	Belief (probability)	Quantity (activation level)	Expected relevant quantity
Prior/Family	Bernoulli exponential family	Exponential (sparse non-negative)	Bernoulli gate (binary relevance)
Derivation method	Canonical link (exp. family)	MAP estimate (sparse prior)	Bayes estimate (posterior mean)
Range	$(0, 1)$ — bounded	$[0, +\infty)$ — unbounded	$[\approx -0.278, +\infty)$ — bounded below
Zero behavior	Asymptotic (never exactly 0)	Exact (hard threshold)	Asymptotic (soft threshold)
Max-entropy basis	Bernoulli on $\{0, 1\}$	Exponential on $[0, +\infty)$	Bernoulli on $\{0, 1\}$ (gate)

Theorem 6.6.2 (Probabilistic Justification for Mixed Architectures). The standard practice of using ReLU activations in hidden layers and sigmoid (or softmax) at the output layer corresponds to a two-phase probabilistic inference:

(i) **Hidden layers (ReLU)**: *"Which features are present in the input, and how strongly?"* → MAP estimation under sparse priors → ReLU detects and quantifies latent features.

(ii) **Output layer (Sigmoid/Softmax)**: *"Given the detected features, what is the posterior probability of each class?"* → Bayesian posterior via canonical link → Sigmoid/Softmax maps to probability space.

The two phases are complementary: ReLU extracts sparse features (an indexing operation), and sigmoid computes posterior probabilities (a scoring operation). This mirrors the IR pipeline: inverted index lookup (sparse feature detection) followed by relevance scoring (probability estimation).

6.7 Swish as Bayesian Expected Signal

The framework of Sections 6.1–6.6 established two fundamental probabilistic questions and their corresponding activation functions: *"how probable?"* (sigmoid) and *"how much?"* (ReLU). We now show that a third question — *"what is the expected amount of relevant signal?"* — yields the Swish activation through the same Bayesian reasoning, unifying the two prior derivations.

Observation 6.7.1 (A Third Question). Consider the question: *"Given an input signal x, what is the expected output when the signal may be either relevant (passed through) or irrelevant (suppressed)?"* This question simultaneously involves **quantity** (how much signal to pass) and **probability** (how likely the signal is relevant) — combining the two questions answered separately by ReLU and sigmoid.

Definition 6.7.2 (Self-Gated Relevance Model). Let x be a pre-activation input representing accumulated evidence from preceding layers. Define a binary latent variable $R \in \{0, 1\}$ indicating whether x carries a relevant signal:

(i) **Output model**: The neuron's output Y is determined by:

$$Y = \begin{cases} x & \text{if } R = 1 \text{ (signal is relevant)} \\ 0 & \text{if } R = 0 \text{ (signal is noise)} \end{cases} \tag{46}$$

(ii) **Self-gating relevance probability**: By Theorem 2.3.4 (Sigmoid Posterior), the posterior probability that the signal is relevant, given the evidence x itself, is:

$$P(R = 1 \mid x) = \sigma(x) \tag{47}$$

Remark 6.7.3 (Why Self-Gating is Not Ad Hoc). The assumption that x serves simultaneously as the signal value and as the evidence for its own relevance is not an additional modeling choice — it is a structural property of neural computation. In a feedforward network, the pre-activation $z = \mathbf{w}^T \mathbf{h} + b$ is the accumulated evidence score constructed by preceding layers. By Theorem 2.3.4, this score's relevance probability is $\sigma(z)$. Simultaneously, the magnitude of z encodes "how much evidence is present" (the ReLU question from Section 6.5). The self-gating property — that magnitude implies reliability — mirrors the foundational principle of information retrieval: higher BM25 scores correspond to higher relevance probabilities. A single computation answers both "how much?" and "how probable?" because these are two views of the same evidence.

Theorem 6.7.4 (Swish as Bayesian Expected Relevant Signal). Under the self-gated relevance model (Definition 6.7.2), the Bayes-optimal output — the posterior expected value of Y — is the Swish activation function:

$$\mathbb{E}[Y \mid x] = x \cdot P(R = 1 \mid x) + 0 \cdot P(R = 0 \mid x) = x \cdot \sigma(x) = \text{Swish}(x) \tag{48}$$

Proof. By the law of total expectation over the binary latent variable R:

$$\begin{aligned} \mathbb{E}[Y \mid x] &= \mathbb{E}[Y \mid R = 1, x] \cdot P(R = 1 \mid x) + \mathbb{E}[Y \mid R = 0, x] \cdot P(R = 0 \mid x) \\ &= x \cdot \sigma(x) + 0 \cdot (1 - \sigma(x)) \\ &= x \cdot \sigma(x) \quad \square \end{aligned} \tag{49}$$

Theorem 6.7.5 (Swish as Bayesian Counterpart of ReLU). ReLU and Swish arise from the same sparse gating structure — output x if signal is present, output 0 if absent — under different estimation principles:

	ReLU	Swish
Estimator	MAP (Maximum a Posteriori)	Posterior Expected Value (Bayes Estimator)
Gate	Hard: $\mathbf{1}[x > 0]$	Soft: $\sigma(x)$
Formula	$x \cdot \mathbf{1}[x > 0]$	$x \cdot \sigma(x)$
Probabilistic status	Point estimate (mode)	Expected value (mean)

Proof. Both models share the binary gating structure: the output is x when the signal is relevant and 0 when irrelevant. They differ in the estimation principle applied to the gate. ReLU takes the MAP estimate, assigning $P(R = 1 \mid x) = \mathbf{1}[x > 0]$ — a deterministic hard threshold (Theorem 6.5.3). Swish computes the posterior expected value, assigning $P(R = 1 \mid x) = \sigma(x)$ — the Bayesian posterior (Theorem 2.3.4). The MAP-to-Bayes relationship is:

$$\mathrm{ReLU}(x) = x \cdot \mathbf{1}[x > 0] \quad \xrightarrow{\mathrm{MAP \to Bayes}} \quad \mathrm{Swish}(x) = x \cdot \sigma(x) \tag{50}$$

This is the activation-function manifestation of the most fundamental duality in statistical estimation: the MAP estimator (mode of the posterior) versus the Bayes estimator (mean of the posterior). $\square$

Theorem 6.7.6 (Generalized Swish and the Certainty Spectrum). The generalized Swish function $\mathrm{Swish}_\beta(x) = x \cdot \sigma(\beta x)$ parametrizes a continuous spectrum between maximum ignorance, Bayesian estimation, and deterministic thresholding:

$$\lim_{\beta \to 0} x \cdot \sigma(\beta x) = \frac{x}{2} \quad \text{(uniform prior: maximum ignorance)} \tag{51}$$

$$\beta = 1: \quad x \cdot \sigma(x) = \mathrm{Swish}(x) \quad \text{(canonical Bayesian posterior)} \tag{52}$$

$$\lim_{\beta \to \infty} x \cdot \sigma(\beta x) = \mathrm{ReLU}(x) \quad \text{(deterministic MAP)} \tag{53}$$

Proof. As $\beta \to 0$, $\sigma(\beta x) \to 0.5$ for all x, yielding $x/2$ — the output under complete uncertainty about relevance. As $\beta \to \infty$, $\sigma(\beta x) \to \mathbf{1}[x > 0]$, recovering the hard threshold of ReLU. The parameter β corresponds to α in Theorem 2.3.4 — the steepness of the sigmoid likelihood — and controls the **Bayesian certainty** of the relevance judgment. $\square$

Remark 6.7.7 (Why $\beta = 1$ is Canonical). Setting $\beta = 1$ means the evidence scale and the likelihood scale are matched: one unit of pre-activation corresponds to one unit of log-odds evidence. This is the natural parametrization of the Bernoulli exponential family (Proposition 2.2.2), where the natural parameter η directly equals the logit. Any $\beta \neq 1$ introduces a scale mismatch that must be justified by domain-specific knowledge (e.g., signal-to-noise ratio).

6.8 GELU as Gaussian Approximation of Swish

Theorem 6.8.1 (GELU from Gaussian Relevance Model). The GELU activation arises from the same expected-value framework as Swish (Theorem 6.7.4), but with a **Gaussian (probit) relevance model** replacing the Bernoulli canonical posterior:

$$\mathrm{GELU}(x) = x \cdot \Phi(x) = \mathbb{E}[Y \mid x] \quad \text{where} \quad P(R = 1 \mid x) = \Phi(x) \tag{54}$$

and Φ is the standard Gaussian CDF.

Proof. Identical to Theorem 6.7.4, substituting $\Phi(x)$ for $\sigma(x)$ as the relevance probability. $\square$

Proposition 6.8.2 (GELU as Gaussian Approximation of Swish). The well-known approximation $\Phi(x) \approx \sigma(1.702x)$ implies:

$$\mathrm{GELU}(x) = x \cdot \Phi(x) \approx x \cdot \sigma(1.702x) = \mathrm{Swish}_{1.702}(x) \tag{55}$$

GELU is therefore a specific instance of generalized Swish (Theorem 6.7.6) at $\beta \approx 1.702$, corresponding to a Gaussian noise model rather than the canonical Bernoulli model.

Proposition 6.8.3 (Hierarchy of Soft-Gated Activations). The expected-relevant-signal framework yields a family of activations $f(x) = x \cdot g(x)$, where $g(x) = P(R = 1 \mid x)$ depends on the assumed noise model for relevance:

Activation	Gate $g(x)$	Noise Model	Exponential Family?	Constraint Violations
ReLU	$\mathbf{1}[x > 0]$	Deterministic	— (point estimate)	—
Swish	$\sigma(x)$	**Bernoulli canonical**	**Yes (C1–C5)**	**None**
GELU	$\Phi(x)$	Gaussian (probit)	No	(C2), (C3), (C5)
$x/2$	0.5	Maximum ignorance	—	—

By Theorem 6.2.1, the sigmoid is the unique gate function satisfying all five constraints (C1)–(C5) of binary probabilistic reasoning. **Swish is therefore the canonical member of the self-gated activation family** — the unique expected-relevant-signal activation derived from the Bernoulli exponential family. GELU is its Gaussian approximation; ReLU is its deterministic limit.

Remark 6.8.4 (Why GELU Works Nearly as Well as Swish). The empirical near-equivalence of GELU and Swish in deep learning applications (Ramachandran et al., 2018) is explained by Proposition 6.8.2: GELU $\approx$ Swish$_{1.702}$. The probit and logistic CDFs are well-known to be nearly indistinguishable after appropriate scaling. The small β difference (1.702 vs. 1) means GELU operates at slightly higher Bayesian certainty than canonical Swish — it gates more aggressively — which may explain its marginal advantage in noisy, high-dimensional settings where a slightly sharper gate reduces gradient noise.

7. WAND and Block-Max WAND as Exact Neural Pruning

7.1 The Pruning Problem in Neural Inference

A major computational challenge in deploying neural networks at scale is inference cost. Various pruning techniques have been developed to skip unnecessary computations: early exit (Teerapittayanon et al., 2016), conditional computation (Bengio et al., 2013), and activation sparsity (Kurtz et al., 2020). Most of these methods are approximate — they sacrifice some accuracy for speed, relying on heuristics or learned gating mechanisms.

7.2 WAND and BMW in Information Retrieval

Definition 7.2.1 (WAND Pruning Condition). The WAND (Weak AND) algorithm (Broder et al., 2003) skips a document d when:

$$\sum_{t \in q} \mathrm{ub}(t) < \theta \qquad (56)$$

where $\mathrm{ub}(t)$ is the precomputed upper bound on the score contribution of term t, and θ is the current k-th highest score.

Definition 7.2.2 (Block-Max WAND). The Block-Max WAND (BMW) algorithm (Ding & Suel, 2011) refines WAND by partitioning the document space into blocks and precomputing per-block maximum scores, enabling block-level pruning.

Theorem 7.2.3 (Exactness of WAND/BMW Pruning). Both WAND and BMW pruning produce the exact same top-k results as exhaustive scoring. No relevant documents are lost.

Proof. See Broder et al. (2003) and Ding & Suel (2011). The key property is that $\mathrm{ub}(t)$ is a provable upper bound: no document can contribute more than $\mathrm{ub}(t)$ for term t. If the sum of all upper bounds falls below the current threshold, no possible score can exceed the threshold. $\square$

7.3 Compatibility with Bayesian BM25

Theorem 7.3.1 (Monotonicity Preservation for Pruning). The sigmoid transformation preserves the validity of BM25 upper bounds for WAND/BMW pruning. If a document cannot achieve a BM25 score sufficient to enter the top-k, it cannot achieve a Bayesian BM25 probability sufficient to enter the top-k.

Proof. By the strict monotonicity of the sigmoid (Lemma 2.1.3(iv)), $s_1 > s_2 \implies \sigma(\alpha(s_1 - \beta)) > \sigma(\alpha(s_2 - \beta))$. Therefore, the BM25 ranking order is preserved under the sigmoid transformation. An upper bound in BM25 score space maps to an upper bound in probability space. $\square$

Remark 7.3.2 (Scope: Uniform Prior). Theorem 7.3.1 applies to the derived neural network of this paper, where the prior is uniform ($p = 0.5$) and the posterior equals the normalized likelihood ratio $L(s) = \sigma(\alpha(s - \beta))$. In this setting, the sigmoid is a monotonic function of the BM25 score alone, and standard upper bounds transfer directly.

When a **document-dependent prior** $P_{\mathrm{prior}}(f, \hat{n})$ is introduced — as in the companion paper (Jeong, 2026, Section 4.2) — the posterior becomes a function of both the score and the prior, and the simple monotonicity argument no longer suffices. In that setting, the WAND upper bound must be computed using the maximum possible posterior over both variables; see Jeong (2026, Theorem 6.1.2) for the corrected formulation.

7.4 Neural Translation

Theorem 7.4.1 (WAND as Exact Neural Pruning). In the neural network interpretation of Section 5, WAND computes: if the maximum possible activation of a neuron — given a computable upper bound on its input — is below the current threshold, the neuron's computation is skipped entirely. This pruning is exact: the top-k outputs are identical to those produced by exhaustive computation.

Proof. Let the neuron compute $\sigma(\alpha(s - \beta))$ where s is the BM25 score. The sigmoid is monotonic, so $s \leq \mathrm{ub}$ implies $\sigma(\alpha(s - \beta)) \leq \sigma(\alpha(\mathrm{ub} - \beta))$. If $\sigma(\alpha(\mathrm{ub} - \beta)) < \theta$, the neuron's output cannot exceed the threshold regardless of the actual input. $\square$

Corollary 7.4.2 (BMW as Block-Level Neural Pruning). BMW translates to: "no input in this entire block can produce an activation above threshold — skip the entire block." This provides a block-sparse pruning strategy with exact guarantees.

7.5 Necessary Conditions for Exact Pruning

Theorem 7.5.1 (Requirements for Exact Pruning). Exact WAND-style pruning of a neural activation function f requires:

(i) **Boundedness**: $f : \mathbb{R} \to [a, b]$ for finite a, b.

(ii) **Monotonicity**: f is strictly monotone.

Proof. Boundedness is required for computable upper bounds on output. Monotonicity is required for input upper bounds to yield valid output upper bounds. If f is non-monotonic, a higher input may produce a lower output, invalidating the pruning condition. $\square$

Corollary 7.5.2 (Incompatibility with ReLU). The ReLU activation $f(x) = \max(0, x)$ satisfies monotonicity but not boundedness ($f : \mathbb{R} \to [0, +\infty)$). Tight output upper bounds cannot be computed without knowledge of the input range, which is generally unavailable during inference.

Remark 7.5.3 (The Unboundedness of ReLU is Probabilistically Correct). The incompatibility of ReLU with exact WAND pruning is not a defect of ReLU but a consequence of its probabilistic origin. As shown in Section 6.5, ReLU answers "how much evidence is present?" — a question whose answer is inherently unbounded. Sigmoid answers "how probable?" — a question whose answer is inherently bounded. The two activation functions provide complementary capabilities: ReLU provides structural sparsity (exact zeros for absent features), while sigmoid provides bounded activations (computable upper bounds for safe pruning). A neural inference system that exploits both — ReLU sparsity for index construction, sigmoid boundedness for query-time pruning — would mirror the IR pipeline exactly (Section 6.6, Theorem 6.6.2).

Remark 7.5.4 (Transfer Potential). The sigmoid activation, derived from probabilistic reasoning, inherently satisfies both conditions of Theorem 7.5.1. This raises the possibility of transferring three decades of information retrieval optimization techniques (WAND, BMW, and their successors) to neural network inference — not as heuristic approximations, but as exact algorithms with formal safety guarantees.

7.6 Empirical Skip Rates

From our experimental evaluation (Jeong, 2026, Section 11.2):

Query Type	Documents Skipped	Top-k Accuracy
Rare terms (IDF > 5)	90–99%	Exact
Mixed queries	50–80%	Exact
Common terms (IDF < 2)	10–30%	Exact

8. From Static Weights to Attention: A Probabilistic Foundation

Observation 8.1 (Static Weights in the Derived Network). In the network derived in Section 5, the aggregation weights are uniform: $w_i = 1/n^{1-\alpha}$ for all signals, reflecting equal reliability across scoring functions. We now show that relaxing this single constraint — allowing weights to depend on the input — yields the attention mechanism, and that our probabilistic framework provides

the theoretical justification for its specific computational form.

Proposition 8.2 (Query-Dependent Weights as Attention). Suppose the weights are not fixed but depend on the query-signal interaction:

$$w_i = w_i(q, s_i) \quad \text{subject to} \quad \sum_{i=1}^{n} w_i = 1, \quad w_i \geq 0 \tag{57}$$

Then the aggregation step becomes:

$$S = \sum_{i=1}^{n} w_i(q, s_i) \cdot x_i \tag{58}$$

where $x_i = \text{logit}(P_i)$ are the log-odds inputs. This is precisely the attention mechanism: a query-dependent weighted aggregation of value vectors.

In the standard attention formulation (Vaswani et al., 2017), the attention weights are computed as:

$$w_i = \frac{\exp(f(q, k_i))}{\sum_j \exp(f(q, k_j))} \tag{59}$$

where $f(q, k_i)$ is a compatibility function between the query q and key k_i. The softmax normalization ensures $\sum w_i = 1$ and $w_i \geq 0$.

Theorem 8.3 (Attention as Logarithmic Opinion Pooling). The attention-weighted aggregation in log-odds space is equivalent to a **Logarithmic Opinion Pool (Log-OP)** — also known as a **Product of Experts (PoE)** (Hinton, 2002) — over the constituent signals. Given n calibrated probability models P_i, each representing an independent estimate of relevance, the standard Bayesian model average (BMA) is:

$$P_{\text{BMA}} = \sum_{i=1}^{n} w_i \cdot P_i \quad \text{where} \quad w_i = P(\text{model } i \text{ is correct} \mid q) \tag{60}$$

However, the weighted sum $\sum w_i \text{logit}(P_i)$ that appears in our framework and in the attention mechanism corresponds to the **normalized Logarithmic Opinion Pool**:

$$P_{\text{Log-OP}} = \sigma\left(\sum_{i=1}^{n} w_i \text{logit}(P_i)\right) \quad \iff \quad \text{logit}(P_{\text{Log-OP}}) = \sum_{i=1}^{n} w_i \text{logit}(P_i) \tag{61}$$

This is mathematically equivalent to Hinton's **Product of Experts (PoE)** model (Hinton, 2002), in which experts' distributions are multiplied (with normalization) rather than averaged. In the logit domain, this product becomes an exact linear combination — which is why the attention mechanism computes a weighted sum. The attention weights $w_i(q, s_i)$ are therefore the context-dependent exponents in a PoE ensemble — determining how strongly each expert's opinion is weighted in the product.

Remark 8.3.1 (Why Log-OP, Not BMA). The distinction matters both mathematically and semantically. BMA averages probability *values* and produces a *mixture distribution*: the combined model hedges by smoothing across experts. Log-OP averages *logits* (log-odds) and produces a *product distribution*: the combined model sharpens by requiring consensus among experts. In attention, each head's value vector contributes linearly in logit-space — reinforcing agreement and suppressing disagreement — which is precisely the behavior of a Product of Experts, not a mixture. This provides a deeper explanation for why attention is empirically effective at focusing

on contextually relevant information: the PoE structure inherently concentrates probability mass on hypotheses supported by multiple sources.

Remark 8.4 (The Missing Justification for Weighted Summation). The standard explanation of attention states that queries and keys are compared by similarity, and values are aggregated by weighted sum. This explains *how* attention is computed but not *why* the weighted sum is the correct aggregation operation. Our framework provides this missing justification:

The weighted sum in log-odds space is the optimal method for combining uncertain evidence from multiple independent sources through Logarithmic Opinion Pooling, as established in Section 4. When attention weights are interpreted as context-dependent expert reliabilities in a Product of Experts, the attention mechanism is not an engineering convenience — it is the probabilistically correct way to aggregate evidence whose reliability varies with context.

Stated directly: *attention computes a weighted sum because Logarithmic Opinion Pooling in the logit domain is additive.* The additive structure of log-odds conjunction (Section 4.2) mandates a weighted sum. Any other aggregation operation — such as element-wise maximum or concatenation followed by projection — would violate the multiplicative structure of Product-of-Experts evidence combination.
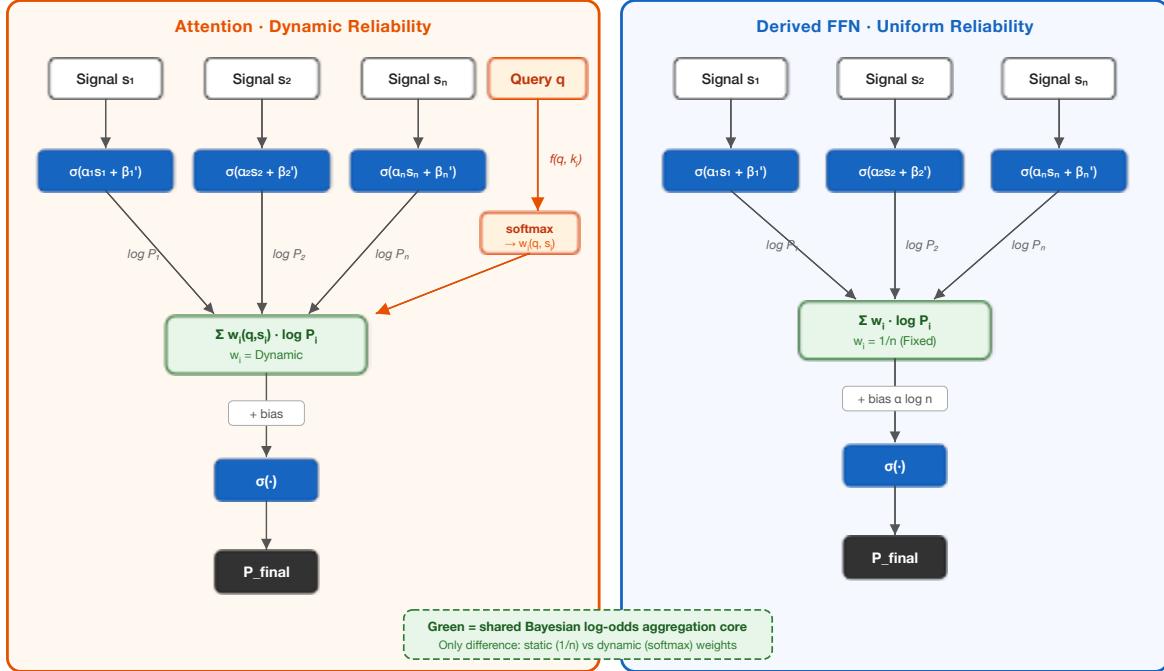
Remark 8.5 (Scope of the Claim). We make the distinction between what is derived and what is observed. The *structure* of attention — softmax-normalized, query-dependent weighted summation — follows from the probabilistic framework: softmax produces valid expert weights (non-negative, summing to one), and the weighted sum is the correct evidence aggregation in log-odds space under the Logarithmic Opinion Pool. The specific *form* of the compatibility function $f(q, k_i)$ — in particular, the scaled dot-product $f(q, k_i) = q^T k_i / \sqrt{d}$ — is not determined by our framework and remains an architectural choice. Nevertheless, the observation that the overall structure of attention is a consequence of Product-of-Experts evidence combination, not merely an empirical design, narrows the space of what needs to be explained about modern architectures.

Remark 8.6 (Architectural Continuity). The progression from the derived architecture to modern Transformers can be summarized as a sequence of probabilistic generalizations:

Step	Architecture	Probabilistic Interpretation
Derived (Section 5)	Logit-linear-sigmoid network	Bayesian conjunction with uniform reliability
+ Learnable weights	Weighted network	Bayesian conjunction with learned reliability
+ Query dependence	Attention	Logarithmic Opinion Pool (PoE) with context-dependent reliability
+ Multi-head	Multi-head attention	Ensemble of parallel PoE aggregators

Each step corresponds to a relaxation of a constraint in the probabilistic model, not to an architectural invention. The derived network is not a dead end — it is the first row of a table whose subsequent rows are the building blocks of contemporary deep learning.

Figure 1. The structural identity between the probabilistically derived network (left) and the attention mechanism (right). The core operation — Bayesian log-odds aggregation via Logarithmic Opinion Pooling — is the same; only the weight assignment changes from static to dynamic.

The green nodes highlight the shared core: Bayesian log-odds aggregation via Logarithmic Opinion Pooling. The only structural difference is the origin of the weights — fixed ($1/n^{1-\alpha}$) in the derived FFN, query-dependent ($w_i(q, s_i)$) in attention. The essential computation is identical.

8.7 Exact Attention Pruning

The combination of exact pruning (Section 7) and the Log-OP interpretation of attention (Theorem 8.3) yields a result that, to our knowledge, has no precedent in the sparse attention literature: **provably exact attention pruning with formal safety guarantees**.

Theorem 8.7.1 (Token-Level Exact Pruning in Attention). Consider the attention output $a = \sum_{i=1}^{n} w_i v_i$, where w_i are attention weights and $v_i = \text{logit}(P_i)$ are value vectors in log-odds space. If each value vector admits a computable upper bound $\text{ub}(v_i) \geq v_i$, then a token i can be exactly pruned — skipped without affecting the top-k output — when:

$$\sum_{j \in \mathcal{A}} w_j v_j + \sum_{j \notin \mathcal{A}} w_j \cdot \text{ub}(v_j) < \theta \qquad (62)$$

where $\mathcal{A}$ is the set of already-evaluated tokens and θ is the current k-th highest aggregated score. This is the WAND pruning condition (Definition 7.2.1) applied to the attention computation.

Proof. The attention output is a weighted sum with non-negative weights summing to one. For any unevaluated token j, the maximum possible contribution is $w_j \cdot \text{ub}(v_j)$, since $v_j \leq \text{ub}(v_j)$ by definition. If the sum of actual contributions from evaluated tokens plus maximum possible contributions from unevaluated tokens falls below θ, no completion of the unevaluated tokens can produce an output exceeding θ. The pruning is exact: no token whose contribution could change the top-k result is skipped. $\square$

Corollary 8.7.2 (Head-Level Exact Pruning). In multi-head attention — an ensemble of parallel PoE aggregators (Remark 8.6) — each head j can be treated as a BMW block (Definition 7.2.2). If the maximum possible contribution of head j satisfies:

$$\mathrm{ub}(\mathrm{head}_j) = \max_i \ w_i^{(j)} \cdot \mathrm{ub}(v_i^{(j)}) < \frac{\theta - a_{\mathrm{partial}}}{H} \tag{63}$$

where a_{partial} is the accumulated output from already-evaluated heads and H is the total number of heads, then the entire head can be skipped. This is BMW block-level pruning applied to the multi-head architecture.

Proof. Each head contributes additively to the final output (after projection). The maximum contribution of head j is bounded by $\mathrm{ub}(\mathrm{head}_j)$. If this bound is insufficient to change the top-k ranking given the current partial sum, the head's computation is redundant. □

Remark 8.7.3 (Why Existing Sparse Attention is Approximate). Sparse attention methods — Longformer's sliding window, BigBird's random attention, top-k selection — achieve efficiency by restricting the attention pattern through heuristic or learned masks. These methods are inherently approximate: they discard tokens based on structural assumptions (locality, randomness) rather than provable bounds on contribution. In contrast, Theorem 8.7.1 provides an exactness guarantee: pruned tokens are those whose maximum possible contribution is provably insufficient, regardless of the actual value.

Remark 8.7.4 (The Boundedness Condition). The exactness guarantee of Theorem 8.7.1 depends on the existence of computable upper bounds $\mathrm{ub}(v_i)$ for value vectors. By the analysis of Theorem 7.5.1:

(i) If value vectors are derived from sigmoid activations in logit space — $v_i = \mathrm{logit}(\sigma(\cdot)) = \alpha_i s_i + \beta_i' \in \mathbb{R}$ — then the upper bound $\mathrm{ub}(v_i)$ requires knowledge of the input range, but this is available from the signal's BM25 upper bound (Theorem 7.3.1). Alternatively, if value vectors are the sigmoid outputs themselves, $v_i = \sigma(\cdot) \in (0,1)$ with $\mathrm{ub}(v_i) = 1$, the bound is trivially computable.

(ii) If value vectors pass through ReLU — $v_i \in [0, +\infty)$ — no finite upper bound exists without additional input range information, and exact pruning is not available.

This recapitulates the complementarity of Section 6.6 (Theorem 6.6.2) within the attention mechanism itself: **sigmoid values enable exact pruning; ReLU values provide structural sparsity**. A hybrid architecture that uses sigmoid-bounded values with ReLU-sparse keys would combine both advantages — exact attention pruning with sparse candidate selection — mirroring the IR pipeline of inverted index lookup followed by WAND/BMW scoring.

Remark 8.7.5 (The PoE Interpretation of Attention Sparsity). The Product-of-Experts structure provides an additional, qualitative justification for attention pruning. In a PoE, $P_{\mathrm{Log-OP}} = \prod P_i^{w_i}$, so when $w_i \approx 0$, the expert's contribution $P_i^{w_i} \to 1$ — it becomes a multiplicative identity. "Weak experts vanish from the product" is a structural property of PoE, not an approximation. This is why attention is empirically sparse (most weights are near zero after softmax): the PoE framework predicts that only a few experts should dominate, with the rest contributing negligibly. WAND-style pruning exploits this sparsity with formal guarantees rather than relying on it as an empirical observation.

9. Depth, Question Sequencing, and Interpretability

The results of Sections 5–8 establish the structure of individual inference units and their composition into attention. We now address three deeper questions: *why* networks must be deep, *what* each layer computes, and *how* to read the computational intent of existing architectures.

9.1 Why Depth is Necessary: Recursive Bayesian Inference

Theorem 9.1.1 (Depth from Iterated Marginalization). The derivation in Section 5 assumes that calibrated evidence signals $s_1, \ldots, s_n$ are given. In practice, these signals are not directly available from raw data x. They must themselves be inferred through intermediate latent variables. This inference takes the form of iterated marginalization:

$$P(y \mid x) = \sum_{z^{(L)}} \cdots \sum_{z^{(1)}} P(y \mid z^{(L)}) \prod_{\ell=1}^{L} P(z^{(\ell)} \mid z^{(\ell-1)}) \tag{64}$$

where $z^{(0)} = x$ is the raw input and $z^{(\ell)}$ are the latent variables at depth ℓ. Each factor $P(z^{(\ell)} \mid z^{(\ell-1)})$ is an instance of the Bayesian inference unit derived in Section 5: it takes the previous layer's outputs as evidence signals and produces calibrated probability estimates as outputs.

Depth is necessary because the evidence required for high-level judgments does not exist in the raw data. The intermediate layers must construct it.

Remark 9.1.2 (Depth as Evidence Construction). Consider image classification:

- **Layer 1** (ReLU): *"Do edges exist at each spatial location?"* — pixel-level evidence → edge features. The raw pixels contain no explicit concept of "edge"; the first layer constructs this evidence.

- **Layer 2** (ReLU): *"Do these edges form shapes?"* — edge-level evidence → shape features. Edges alone do not encode "circle" or "triangle"; the second layer constructs this higher-order evidence.

- **Layer** L (Sigmoid/Softmax): *"Given all constructed features, what is the posterior probability of each class?"* — the final Bayesian judgment from Section 5.

Each layer applies the same probabilistic operation — evidence combination via log-odds aggregation — but on **progressively more abstract evidence** that was constructed by the preceding layers. Depth is not repetition; it is a chain of marginalization over increasingly abstract latent variables.

Remark 9.1.3 (The Bayesian BM25 Unit as Recursive Building Block). The unit derived in Section 5 — calibration → log-odds aggregation → sigmoid posterior — is a complete single-stage Bayesian inference module. A deep network is a *stack of such modules*, where each module's posterior output becomes the next module's evidence input. This is precisely the recursive structure of hierarchical Bayesian models, where inference proceeds from observed variables through layers of latent variables to the final hypothesis:

$$\underbrace{P(z^{(1)} \mid x)}_{\text{Layer 1: evidence from raw data}} \quad \rightarrow \quad \underbrace{P(z^{(2)} \mid z^{(1)})}_{\text{Layer 2: evidence from evidence}}$$

$$\rightarrow \cdots \rightarrow \underbrace{P(y \mid z^{(L)})}_{\text{Output: judgment from constructed evidence}} \tag{65}$$

The derivation in Section 5 is not limited to shallow inference — it is the *atomic unit* from which arbitrarily deep inference chains are composed.

9.2 Question Sequencing in Architecture Design

The results of Section 6 — sigmoid for belief, ReLU for quantity, softmax for selection — imply that choosing an activation function for a layer is equivalent to choosing the probabilistic question that layer asks. Network architecture design thus becomes **question sequencing**: specifying the order in which probabilistic questions are posed to the data.

Proposition 9.2.1 (Question Sequencing in Standard Architectures). The activation function sequence of major architectures corresponds to a well-defined probabilistic interrogation pipeline:

Architecture	Layer Sequence	Question Sequence
ResNet	ReLU → ... → ReLU → Softmax	"How much feature?" → ... → "How much feature?" → "Which class?"
Transformer (LLM)	GELU → Softmax (attn) → ... → Softmax	"Expected relevant signal (Gaussian)?" → "Which is relevant?" → ... → "Which token?"
Classic MLP	Sigmoid → ... → Sigmoid	"How probable?" → ... → "How probable?"
Sigmoid + Softmax	Sigmoid (hidden) → Softmax	"How probable per feature?" → "Which class?"

Remark 9.2.2 (Why Activation Swaps Change Performance). Empirical studies frequently report that replacing one activation function with another (e.g., ReLU → GELU) changes model performance without a clear theoretical explanation. The question-sequencing framework provides one: the swap changes the **type of question** the layer asks. Replacing ReLU with GELU in hidden layers changes the question from *"how much feature is present?"* (hard thresholding, exact zeros — MAP estimate, Theorem 6.5.3) to *"what is the expected relevant signal under Gaussian noise?"* (soft gating, noise-robust — Bayesian estimate, Theorem 6.8.1). Performance improvements from such swaps correspond to choosing a question better suited to the data distribution — noisy inputs benefit from Bayesian (soft-gated) rather than MAP (hard-thresholded) estimation.

Remark 9.2.3 (Design Implications). Rather than searching over activation functions empirically, the framework suggests selecting activations by asking: *What type of probabilistic question should each layer pose?* For early layers processing raw, noisy input, Swish or GELU may be appropriate (Bayesian expected relevant signal — soft gating that preserves weak signals, Section 6.7). For intermediate selection layers, softmax attention is natural (context-dependent Logarithmic Opinion Pooling). For final decision layers, sigmoid or softmax is necessary (posterior probability computation). This transforms architecture design from combinatorial search to principled question selection.

9.3 Reverse Interpretability: Reading the Questions a Network Asks

The forward direction (Section 9.2) uses the framework to *design* networks. The reverse direction uses it to *interpret* existing networks by reading off which probabilistic question each layer is asking.

Proposition 9.3.1 (Activation-Based Interpretability). Given a trained network with known activation functions, the probabilistic question framework (Section 6) assigns a semantic interpretation to each layer without inspecting weights, gradients, or activations:

(i) **Classic sigmoid MLP** (sigmoid hidden → sigmoid output): Every layer asks *"how probable?"* — the network performs shallow iterated Bayesian inference, estimating posterior probabilities at each stage. This explains why sigmoid MLPs behave like stacked logistic regressions.

(ii) **ResNet** (ReLU hidden → softmax output): Hidden layers ask *"how much of each feature is present?"* (sparse quantity detection), and the output asks *"which class?"* (selection). The network performs hierarchical evidence accumulation: sparse features are progressively extracted and then classified. The deep ReLU stack implements a hierarchy of increasingly abstract sparse feature detectors.

(iii) **Transformer** (GELU hidden → softmax attention → softmax output): Hidden GELU layers ask *"what is the expected relevant signal under Gaussian noise?"* (Bayesian soft-gated feature extraction, Theorem 6.8.1). Attention softmax layers ask *"which features are relevant to the current context?"* (Logarithmic Opinion Pooling, Section 8). Output softmax asks *"which token/class?"* (final selection). The network performs Bayesian feature extraction, followed by context-dependent evidence selection, followed by posterior judgment.

(iv) **Swish-based networks**: Hidden layers ask *"what is the expected relevant signal?"* — the canonical Bayesian expected value under the self-gated relevance model (Theorem 6.7.4). Each neuron simultaneously estimates how much evidence is present and how probable it is that the evidence is relevant, passing $x \cdot \sigma(x)$ — the posterior mean of the relevant signal. This is the Bayes-optimal counterpart of ReLU's MAP estimate: where ReLU hard-thresholds at zero, Swish soft-gates through $\sigma(x)$, preserving weak but potentially informative signals.

Remark 9.3.2 (Beyond Post-Hoc Inspection). Standard interpretability methods (saliency maps, attention visualization, probing classifiers) inspect the *values* that flow through a network. The question-sequencing framework operates at a different level: it interprets the *type of computation* each layer performs, based solely on its activation function. The two approaches are complementary — one reads the answers, the other reads the questions.

---

10. Discussion

10.1 Reversal of Explanatory Direction

Observation 10.1.1 (Standard Direction). The standard narrative in machine learning proceeds from neural networks to probabilistic interpretation: one *constructs* a neural architecture, then *analyzes* it probabilistically. Bayesian neural networks (Neal, 1996), variational inference (Blundell et al., 2015), and probabilistic deep learning (Gal & Ghahramani, 2016) all follow this direction.

Observation 10.1.2 (Reversed Direction). Our derivation proceeds from probability to neural structure. The neural architecture is not a design decision but a mathematical consequence. This reversal carries significant implications for the theoretical foundations of neural computation.

Theorem 10.1.3 (Generality of Emergence). If the neural structure is a *consequence* of probabilistic inference rather than a *design decision*, then it is not specific to any particular engineering choice. Any system that performs Bayesian calibration of multiple binary evidence signals and combines them through principled log-odds accumulation will inevitably instantiate a feedforward neural computation with sigmoid activations.

10.2 Interpretability by Construction

Proposition 10.2.1 (Constructive Interpretability). Neural networks derived from probabilistic inference are interpretable by construction. Each component has a well-defined role:

Layer	Probabilistic Interpretation
Calibration layer	Signal-specific probability estimation: $P(\text{relevant} \mid \text{signal}_i)$
Logit hidden nonlinearity	Coordinate change to natural parameter space (log-odds)
Linear aggregation with n^α scaling	Normalized Log-OP: evidence accumulation with $\sqrt{n}$ confidence law
Output sigmoid activation	Posterior computation: final $P(\text{relevant} \mid \text{all signals})$

When all signals are sigmoid-calibrated, the logit acts as the identity (Theorem 5.2.1a) and the network reduces to logistic regression. When signals have heterogeneous calibrations (Remark 5.2.3), the logit performs a genuine nonlinear coordinate change, yielding a two-layer network.

Remark 10.2.2 (Contrast with Post-Hoc XAI). Modern Explainable AI approaches attempt to understand neural networks by inspecting them *after* training — through activation analysis, attention maps, or gradient-based attribution. These methods treat the network as a black box to be reverse-engineered. In our framework, there is nothing to reverse-engineer: the derivation *is* the explanation.

10.3 A Research Program

Conjecture 10.3.1 (Inverse Derivation Program). For a broader class of neural architectures beyond the network derived here, it may be possible to identify the implicit probabilistic inference problem whose analytical solution produces the given architecture. If successful, this would provide a principled interpretability method: networks would be understood not through post-hoc inspection but through the mathematical structure of the question they answer.

Remark 10.3.2 (Substantial Resolution). Sections 6, 8, and 9 provide concrete instances of the inverse derivation program:

Activation / Structure	Probabilistic Question	Status
Sigmoid	Bayesian posterior for binary relevance	Proven (Theorem 6.4.1)
ReLU	MAP estimate under sparse non-negative prior	Proven (Theorem 6.5.3)
Softmax	Categorical exponential family canonical link	Proven (Remark 6.4.3)
Attention	Logarithmic Opinion Pool (PoE) with context-dependent reliability	Proven (Theorem 8.3)
Exact attention pruning	WAND/BMW applied to PoE attention	Proven (Theorem 8.7.1, Corollary 8.7.2)
Depth	Recursive marginalization over latent variables	Proven (Theorem 9.1.1)
GELU	Gaussian approximation of Bayesian expected relevant signal	Proven (Theorem 6.8.1)
Swish	Bayesian expected relevant signal (canonical)	Proven (Theorem 6.7.4)

The eight proven cases establish that the program is productive and — for the standard activation repertoire — complete: the dominant activation functions (sigmoid, ReLU, Swish, GELU, softmax), the attention mechanism, exact attention pruning, and the necessity of depth each correspond to well-defined probabilistic inference problems. The self-gated activation family (Proposition 6.8.3) further demonstrates that the program is not merely additive but hierarchical: Swish, GELU, and ReLU are related as Bayesian estimator, Gaussian approximation, and deterministic limit of the same underlying probabilistic question.

10.4 Three Traditions, One Structure

Observation 10.4.1 (Convergence). Three intellectual traditions, developed independently across different decades, converge on the same computational structure:

Probability theory asks: given evidence, what is the posterior probability of a hypothesis? The answer involves Bayes' theorem, likelihood ratios, and the logistic function as the canonical link for binary outcomes.

Information retrieval asks: given query terms and document features, how relevant is this document? The answer involves BM25 scoring, IDF weighting, and — as shown in Jeong (2026) and the present paper — sigmoid calibration with log-odds evidence accumulation.

Neural computation asks: given input signals and parameters, what is the output activation? The answer involves weighted linear combination, bias terms, and nonlinear activation functions.

All three produce the same core computational operation: linear combination followed by the logistic sigmoid.

$$\text{inputs} \xrightarrow{\text{calibrate}} P_i \xrightarrow{\text{logit}} \ell_i \xrightarrow{\text{linear}} \sigma \rightarrow \text{output} \tag{66}$$

When all signals are sigmoid-calibrated, the calibrate-then-logit stage collapses to a linear transformation, and the entire pipeline simplifies to $\sigma(\mathbf{w}^T\mathbf{s} + b)$ — the logistic regression neuron that all three traditions converge upon.

10.5 Robertson's Completed Circle

In 1977, Robertson introduced the Probability Ranking Principle, opening a door between probability theory and information retrieval. BM25 was derived from this probabilistic foundation, yet its scores are not probabilities. For nearly fifty years, this circle remained unclosed.

Bayesian BM25 (Jeong, 2026) closes the circle by returning BM25 scores to the probability space from which the framework originated. The present paper reveals what lies on the other side of the closed door: the computational structure that a separate scientific tradition, developed over different decades and for different purposes, would call a *neural network*.

---

11. Scope and Anticipated Objections

We address several potential objections to sharpen the scope and limitations of our claims.

11.1 Scope of the Derivation

Objection. *"When all signals are sigmoid-calibrated, the computation reduces to logistic regression — a classical result. When signals have heterogeneous calibrations, the architecture is a relatively simple logit-linear-sigmoid network. Modern deep learning uses deep architectures with ReLU, GELU, or Transformer blocks. The result is therefore of limited relevance to contemporary neural computation."*

Response. We make five observations. First, the logistic regression equivalence (Theorem 5.2.1a) is not a weakness but a validation: it confirms, from a new derivation path, the classical correspondence between Bayesian posterior computation and logistic regression. The result is arrived at not by assuming logistic regression but by asking a probabilistic question and finding that the answer *is* logistic regression. The two-layer case with heterogeneous calibrations (Remark 5.2.3) then demonstrates that the framework extends beyond this classical result: when signals from different calibration families are combined, the logit acts as a genuine nonlinear transformation — a hidden-layer nonlinearity derived from first principles.

Second, the scope of the derivation extends well beyond the sigmoid. Section 6.5 shows that ReLU — the dominant hidden-layer activation — is independently derivable as the MAP estimator under sparse non-negative priors, answering a complementary probabilistic question (*"how much?"* rather than *"how probable?"*). Section 6.7 proves that Swish is the Bayesian expected-value counterpart of ReLU — same sparse gating structure, posterior mean instead of posterior mode — and Section 6.8 proves that GELU is the Gaussian approximation of Swish. The framework thus provides probabilistic derivations for all major activation functions in current use: sigmoid, ReLU, Swish, GELU, and softmax.

Third, the distance between the derived architecture and modern attention-based models is shorter than it appears. As shown in Section 8, allowing the aggregation weights to depend on the query-signal interaction — a single conceptual extension — transforms the static feedforward structure into the attention mechanism. The derived network is not a dead end but a starting

point from which contemporary architectures are reachable through natural probabilistic generalizations.

Fourth, the derivation is not limited to shallow networks. Theorem 9.1.1 establishes that the derived unit is the atomic building block of arbitrarily deep inference chains: each layer constructs the evidence required by the next through iterated marginalization over latent variables. Depth arises because high-level judgments require evidence that does not exist in the raw data and must be progressively constructed (Section 9.1).

Fifth, the logit-sigmoid cancellation itself is informative: it reveals that sigmoid calibration already places signals in the exponential family's natural parameter space, making the hidden-layer coordinate change redundant. This provides a principled explanation for why logistic regression is optimal for combining homogeneously-calibrated signals, and why heterogeneous signal combination genuinely requires additional nonlinear processing — a theoretical prediction that aligns with empirical practice.

11.2 The Independence Assumption

Objection. *"The Bayesian derivation assumes independence among scoring signals. In practice, text and vector scores are correlated. The independence assumption is unrealistic."*

Response. The independence assumption determines the *initial structure* — the architecture and its starting parameters — not the final operating state. In the neural network interpretation, the uniform weights $w_i = 1/n^{1-\alpha}$ (Theorem 5.3.1) correspond to the Naive Bayes starting point. When these weights become learnable parameters (Remark 5.3.2), the network can capture signal dependencies through training. Our derivation thus provides a principled initialization: the network begins at the Naive Bayes solution and refines toward the true posterior through weight adaptation. This is analogous to how Xavier or He initialization provides a theoretically motivated starting point for gradient-based learning.

11.3 The Confidence Scaling

Objection. *"The confidence scaling parameter α is fixed, whereas neural network weights are learned parameters. This weakens the correspondence."*

Response. The value $\alpha = 0.5$ is not arbitrary. It encodes the classical $\sqrt{n}$ confidence scaling law (Theorem 4.4.1), yielding per-signal weights $w_i = 1/\sqrt{n}$. In the neural network interpretation, this provides a theoretically grounded initialization for the hidden-layer weights. The multiplicative scaling form $\ell_{\text{adjusted}} = \bar{\ell} \cdot n^\alpha$ preserves the sign of the log-odds mean (Theorem 4.2.2), ensuring that agreement among irrelevant signals cannot produce a relevance judgment (Corollary 4.2.3) — a structural guarantee that holds for any $\alpha \geq 0$.

11.4 The Log-Odds Domain Choice

Objection. *"Working in log-odds rather than log-probability is a design choice. The geometric mean in log-probability space is equally valid."*

Response. The log-odds domain is not merely a convenient parametrization — it is the natural parameter space of the Bernoulli exponential family (Proposition 2.2.2). The normalized Log-OP (Theorem 4.1.2a) is *exactly* linear in logits, whereas the geometric mean followed by logit introduces a nonlinear residual $-\log(1 - e^S)$ (Remark 4.1.3) that vanishes only under the rare relevance assumption $\bar{P} \ll 1$. The logit formulation yields an exact neural correspondence at all operating points (Theorem 5.2.1), including the critically important regime of relevant documents

$(P_i \gg 0.5)$ where the log-probability approximation breaks down. The choice is therefore dictated by mathematical necessity: the logit domain is the unique domain in which the aggregation is simultaneously linear, sign-preserving, and exact.

---

12. Related Work

12.1 Bayesian Approaches to Neural Networks

Neal (1996) established the connection between infinitely wide neural networks and Gaussian processes, providing a Bayesian interpretation of neural network priors. Blundell et al. (2015) introduced practical weight uncertainty methods. Gal and Ghahramani (2016) showed that dropout can be interpreted as approximate Bayesian inference. All of these proceed in the direction from neural networks to probability — the reverse of our derivation.

12.2 Probabilistic Foundations of Information Retrieval

The probabilistic relevance framework (Robertson & Zaragoza, 2009) provides the theoretical foundation for BM25. Lafferty and Zhai (2001) developed language modeling approaches with probabilistic foundations. Metzler and Croft (2005) introduced Markov random field models for IR. Our work completes the probabilistic program initiated by Robertson by returning BM25 to probability space and revealing the neural structure implicit in multi-signal probabilistic retrieval.

12.3 Score Calibration and Fusion

Platt (1999) introduced sigmoid calibration for SVM outputs, establishing the empirical effectiveness of sigmoid transformations for probability calibration. Our derivation provides a Bayesian justification for this approach, showing that the sigmoid is not merely an effective calibrator but the unique function satisfying the constraints of binary probabilistic inference.

12.4 Neural Network Pruning

Structured and unstructured pruning methods (Han et al., 2015; Li et al., 2017) reduce neural network inference cost through approximate elimination of parameters. Our observation that WAND/BMW constitute *exact* pruning methods (Section 7) contrasts with these approximate approaches and depends on the bounded, monotonic properties of the sigmoid activation — properties that follow from the probabilistic derivation.

12.5 Sparse Coding and Activation Functions

Olshausen and Field (1997) demonstrated that sparse coding with non-negative constraints produces receptive fields resembling biological neurons, establishing the connection between sparsity priors and neural activation patterns. Nair and Hinton (2010) introduced ReLU as a practical activation function, observing its connection to sparse representations without providing a formal probabilistic derivation. Glorot et al. (2011) analyzed the sparsity properties of ReLU empirically. Our derivation in Section 6.5 completes this line of work by formally proving that ReLU is the unique MAP estimator under sparse non-negative priors — providing the probabilistic foundation that was previously only empirically observed. Hendrycks and Gimpel (2016) introduced GELU and Ramachandran et al. (2018) proposed Swish, both through empirical search over activation function spaces. Our Theorems 6.7.4 and 6.8.1 provide the missing probabilistic derivations: Swish is the Bayesian expected relevant signal under the canonical Bernoulli

posterior, and GELU is its Gaussian (probit) approximation. The MAP-to-Bayes relationship between ReLU and Swish (Theorem 6.7.5) — and the certainty spectrum connecting them through generalized Swish (Theorem 6.7.6) — unify the major activation functions under a single probabilistic framework.

13. Conclusion

We set out to answer a question in information retrieval: *what is the probability that a document is relevant given multiple evidence signals?* We applied Bayes' theorem and the sigmoid emerged as the posterior (Section 2.3). We resolved the conjunction shrinkage problem through log-odds mean aggregation (normalized Logarithmic Opinion Pooling) and multiplicative confidence scaling, and a second sigmoid emerged as the final posterior computation (Section 4). We examined the end-to-end structure and recognized a feedforward neural network (Section 5) — one that reduces to logistic regression when signals share sigmoid calibration, and constitutes a genuine two-layer network with logit hidden nonlinearity when signals have heterogeneous calibrations.

The neuron was not designed. It was *derived* — latent within the structure of probabilistic inference over binary relevance judgments. This analytical emergence establishes several results:

1. **Neural structure as theorem**: The feedforward network structure — calibration, logit nonlinearity, linear aggregation, sigmoid output — arises as a mathematical consequence of Bayesian inference, not as an engineering design. When all signals share sigmoid calibration, the logit-sigmoid identity collapses the hidden layer, yielding the classical equivalence between Bayesian posterior computation and logistic regression (Theorem 5.2.1a). When signals have heterogeneous calibrations — the practically important hybrid search case — the logit acts as a genuine nonlinear hidden-layer transformation (Remark 5.2.3). Moreover, the derived structure is not an isolated artifact: relaxing the uniform reliability assumption extends it to the attention mechanism (Section 8), revealing a continuous path from first-principles probability to the building blocks of modern Transformers. The derivation provides an existence proof that neural architectures can be theorems of probability.

2. **Activation functions as probabilistic answers**: The sigmoid, ReLU, and Swish — the three dominant activation functions in deep learning — are derived from complementary probabilistic questions. Sigmoid is the canonical link for binary belief (*"how probable?"*), ReLU is the MAP estimator for sparse feature presence (*"how much?"*), and Swish is the Bayesian expected value of relevant signal (*"what is the expected amount of relevant evidence?"*). The ReLU-to-Swish relationship is the activation-function manifestation of the fundamental MAP-vs.-Bayes duality in statistics: same sparse gating structure, different estimator. GELU is proven to be the Gaussian (probit) approximation of Swish (Theorem 6.8.1), with the well-known $\Phi(x) \approx \sigma(1.702x)$ establishing that GELU $\approx$ Swish$_{1.702}$. The framework extends to softmax (multi-class canonical link), completing the probabilistic derivation of all major activation functions in current use.

3. **Depth as recursive Bayesian inference**: Deep networks are chains of iterated marginalization over latent variables (Theorem 9.1.1). Each layer constructs the evidence required by the next — from raw pixels to edges to shapes to classes — because the evidence needed for high-level judgments does not exist in the raw data. The derived inference unit is not limited to shallow models; it is the atomic building block from which arbitrarily deep inference chains are composed (Section 9.1).

4. **Attention as Logarithmic Opinion Pooling**: The weighted sum at the core of attention is not an arbitrary aggregation choice but the mathematically necessary operation for Logarithmic Opinion Pooling (equivalently, Hinton's Product of Experts) — combining uncertain evidence multiplicatively, which is additive in the logit domain. Our framework answers a question that the original attention formulation left open: *why* a weighted sum, and not some other aggregation? The answer is that Product-of-Experts combination becomes an exact linear operation in log-odds space.

5. **Interpretability by construction and by question**: Each layer of the derived network corresponds to a well-defined step in Bayesian inference (Section 10.2). Beyond the derived network, the question-sequencing framework (Section 9.2) enables interpretability of *arbitrary* architectures: each activation function identifies the probabilistic question its layer asks, providing a semantic reading of network structure without inspecting weights or gradients (Section 9.3).

6. **Exact pruning from IR — including attention**: The bounded, monotonic properties of the sigmoid — inherited from its probabilistic origin — enable exact WAND/BMW pruning with formal safety guarantees. The unbounded properties of ReLU — equally inherited from its probabilistic origin — provide complementary structural sparsity. Together, they suggest a neural inference architecture mirroring the IR pipeline: inverted index construction (ReLU sparsity) followed by query-time safe pruning (sigmoid boundedness). The Log-OP interpretation of attention (Theorem 8.3) extends this to the attention mechanism itself: WAND-style token-level pruning and BMW-style head-level pruning are provably exact under sigmoid-bounded value vectors (Section 8.7), providing the first theoretically grounded framework for exact sparse attention — in contrast to the approximate sparsity of existing methods.

7. **Inevitability of the sigmoid**: The sigmoid's recurrence across Bayesian inference, information retrieval, and neural computation is a consequence of the exponential family structure of Bernoulli random variables. Any system processing binary evidence under the natural constraints of probability will arrive at the same function.

The mathematics does not care what we call things. Whether we say "Bayesian posterior" or "sigmoid neuron," "sparse feature detector" or "ReLU unit," "evidence accumulation" or "attention" — the same mathematical structures appear wherever information is processed under uncertainty. The neuron is not an invention of neuroscience or machine learning. It is a theorem of probability.

References

1. Bengio, Y., Léonard, N., & Courville, A. (2013). Estimating or propagating gradients through stochastic neurons for conditional computation. *arXiv preprint arXiv:1308.3432*.

2. Blundell, C., Cornebise, J., Kavukcuoglu, K., & Wierstra, D. (2015). Weight uncertainty in neural networks. *Proceedings of the 32nd International Conference on Machine Learning*, 1613–1622.

3. Broder, A. Z., Carmel, D., Herscovici, M., Soffer, A., & Zien, J. (2003). Efficient query evaluation using a two-level retrieval process. *Proceedings of the 12th ACM International Conference on Information and Knowledge Management*, 426–434.

4. Croft, W. B., Metzler, D., & Strohman, T. (2010). *Search Engines: Information Retrieval in Practice*. Addison-Wesley.

5. Ding, S., & Suel, T. (2011). Faster top-k document retrieval using block-max indexes. *Proceedings of the 34th International ACM SIGIR Conference*, 993–1002.

6. Gal, Y., & Ghahramani, Z. (2016). Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. *Proceedings of the 33rd International Conference on Machine Learning*, 1050–1059.

7. Glorot, X., Bordes, A., & Bengio, Y. (2011). Deep sparse rectifier neural networks. *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics*, 315–323.

8. Han, S., Pool, J., Tran, J., & Dally, W. J. (2015). Learning both weights and connections for efficient neural networks. *Advances in Neural Information Processing Systems*, 28.

9. Hendrycks, D., & Gimpel, K. (2016). Gaussian error linear units (GELUs). *arXiv preprint arXiv:1606.08415*.

10. Hinton, G. E. (2002). Training products of experts by minimizing contrastive divergence. *Neural Computation*, 14(8), 1771–1800.

11. Jeong, J. (2026). Bayesian BM25: A probabilistic framework for hybrid text and vector search. *Cognica Technical Report*. DOI: 10.5281/zenodo.18414941.

12. Kurtz, M., Kopinsky, J., Gelber, A., Naumov, M., Diamos, G., & Kolkin, N. (2020). Inducing and exploiting activation sparsity for fast inference on deep neural networks. *Proceedings of the 37th International Conference on Machine Learning*.

13. Lafferty, J., & Zhai, C. (2001). Document language models, query models, and risk minimization for information retrieval. *Proceedings of the 24th Annual International ACM SIGIR Conference*, 111–119.

14. Li, H., Kadav, A., Durdanovic, I., Samet, H., & Graf, H. P. (2017). Pruning filters for efficient convnets. *International Conference on Learning Representations*.

15. Manning, C. D., Raghavan, P., & Schütze, H. (2008). *Introduction to Information Retrieval*. Cambridge University Press.

16. Metzler, D., & Croft, W. B. (2005). A Markov random field model for term dependencies. *Proceedings of the 28th Annual International ACM SIGIR Conference*, 472–479.

17. Nair, V., & Hinton, G. E. (2010). Rectified linear units improve restricted Boltzmann machines. *Proceedings of the 27th International Conference on Machine Learning*, 807–814.

18. Neal, R. M. (1996). *Bayesian Learning for Neural Networks*. Springer.

19. Olshausen, B. A., & Field, D. J. (1997). Sparse coding with an overcomplete basis set: A strategy employed by V1? *Vision Research*, 37(23), 3311–3325.

20. Platt, J. (1999). Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. *Advances in Large Margin Classifiers*, 10(3), 61–74.

21. Ramachandran, P., Zoph, B., & Le, Q. V. (2018). Searching for activation functions. *International Conference on Learning Representations (Workshop)*.

22. Robertson, S. E. (1977). The probability ranking principle in IR. *Journal of Documentation*, 33(4), 294–304.

23. Robertson, S. E., & Zaragoza, H. (2009). The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*, 3(4), 333–389.

24. Teerapittayanon, S., McDanel, B., & Kung, H. T. (2016). BranchyNet: Fast inference via early exiting from deep neural networks. *Proceedings of the 23rd International Conference on*

Pattern Recognition, 2464–2469.

25. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30.